



# Machine Learning Project Lifecycle

## Deep Neural Networks Training Focus

Google I/O 2018 Extended  
LaFactory, Technopark, Casablanca, Morocco  
May 12, 2018

Abdelhak Mahmoudi  
[abdelhak.mahmoudi@um5.ac.ma](mailto:abdelhak.mahmoudi@um5.ac.ma)  
 [linkedIn](#)

# Content

- Machine Learning Everywhere
- ML Project Lifecycle
- Deep NN Training Focus
- How can I Learn?
- How can I Apply?
- Demos

# Machine Learning Everywhere

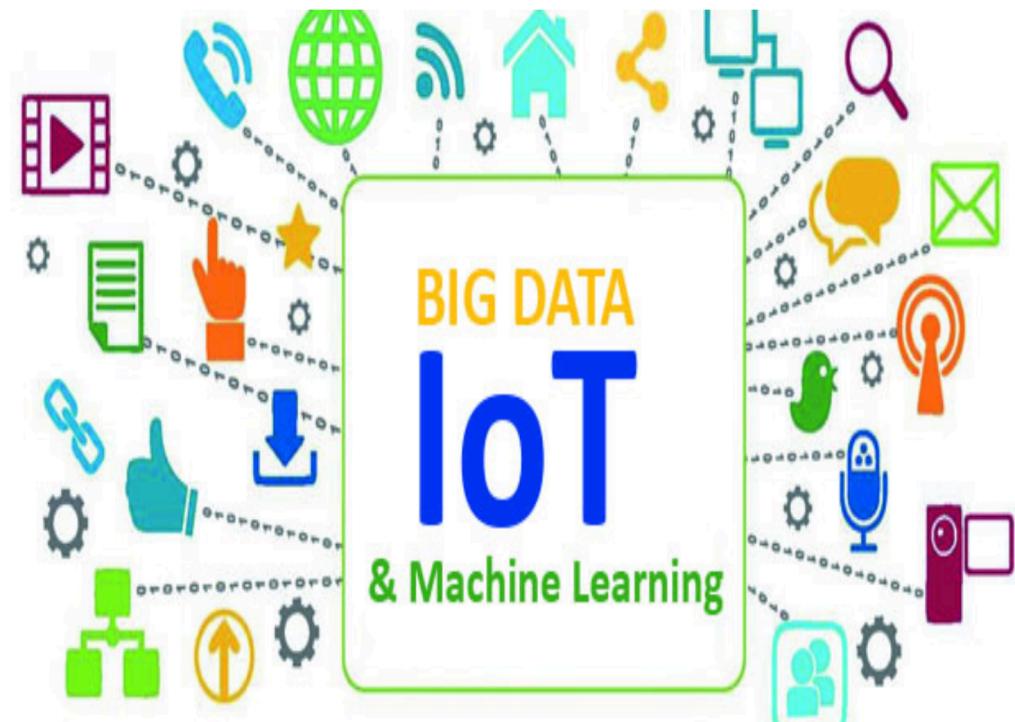
# Use Cases

## **Forbes: The Top 10 AI And Machine Learning Use Cases Everyone Should Know About**

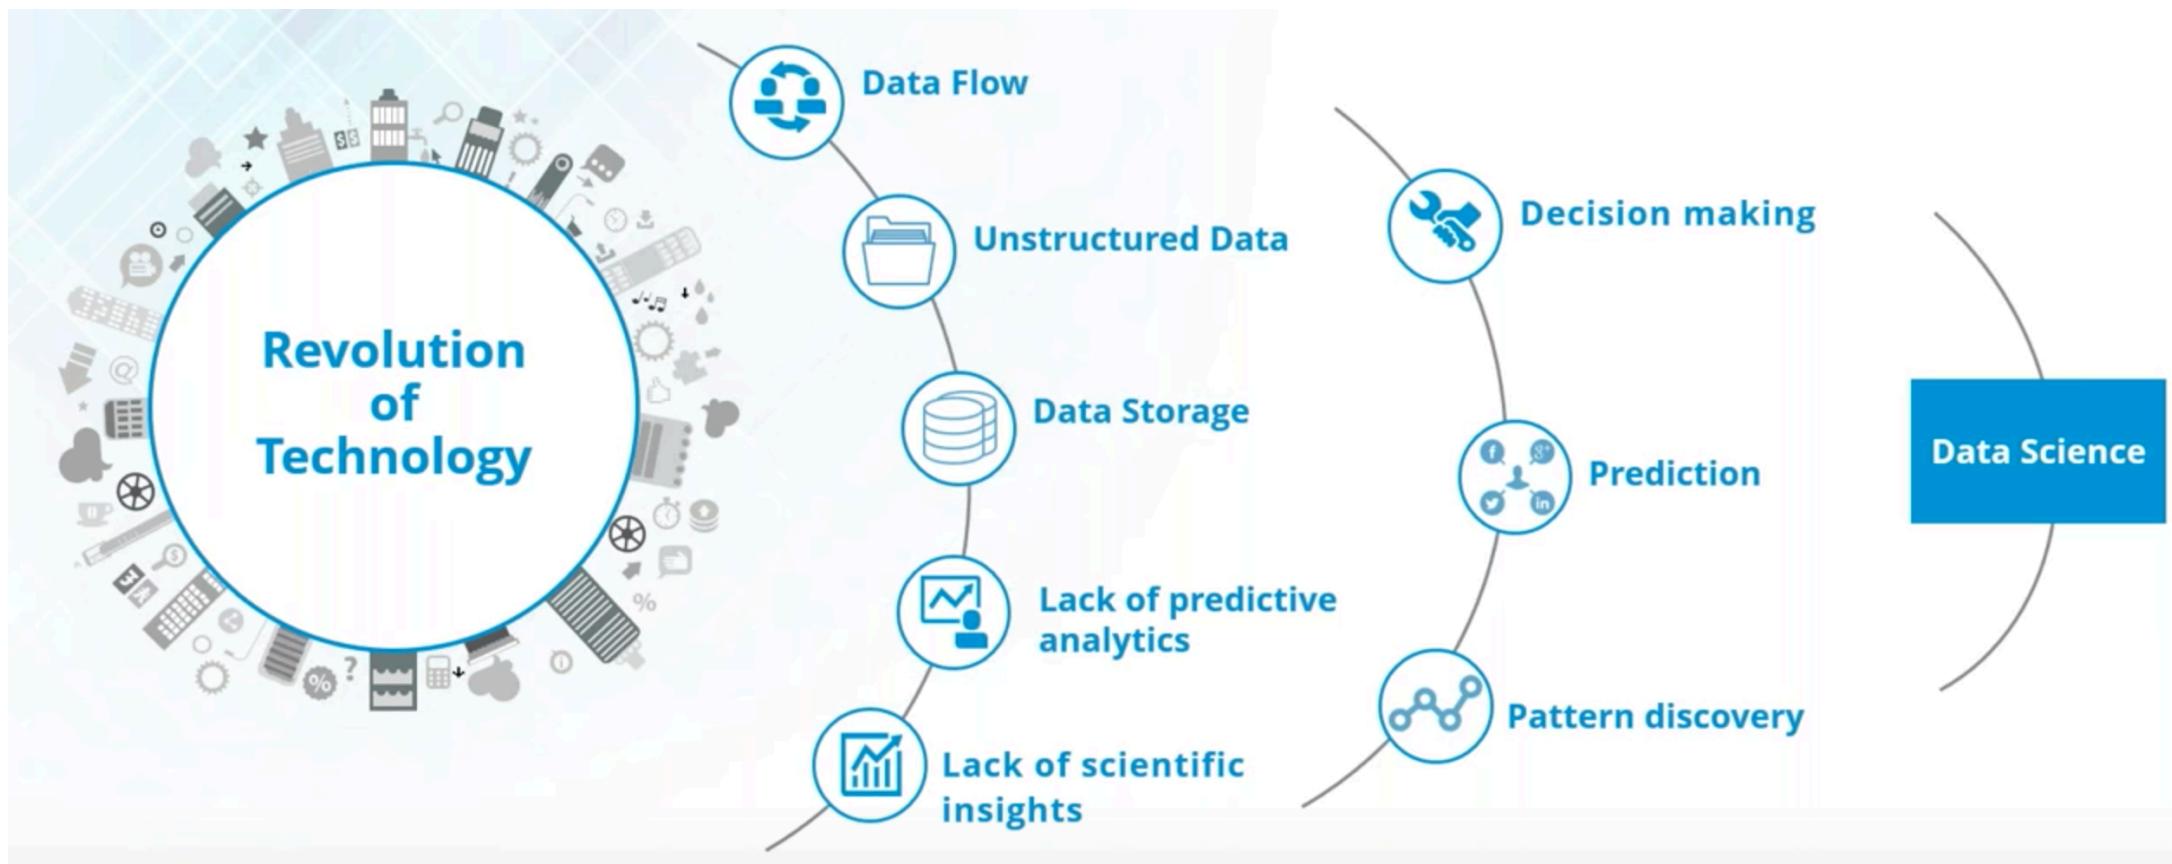
1. Data Security,
2. Personal Security,
3. Financial Trading,
4. Healthcare,
5. Marketing personalization,
6. Fraud Detection,
7. Recommendations,
8. Online Search,
9. NLP,
10. Smart Cars

# Internet of Things

- 25 and 50 billion Internet-connected devices By 2020,
- IoT generates Big Data
  - **Volume**
  - **Velocity** in terms of time and location dependency,
  - **Variety** of multiple modalities
  - **Varying** data quality
- Nowadays, Train in the **Cloud or Fog**, predict locally
- Tomorrow, Train and predict **locally**

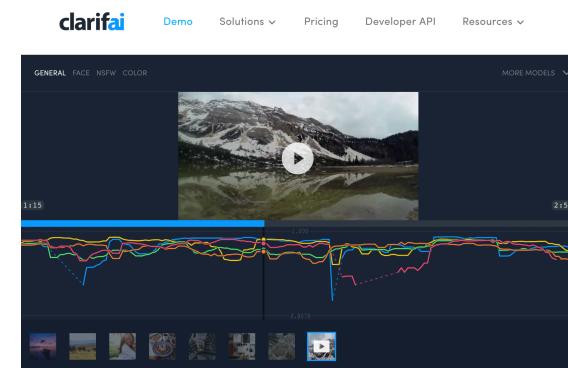
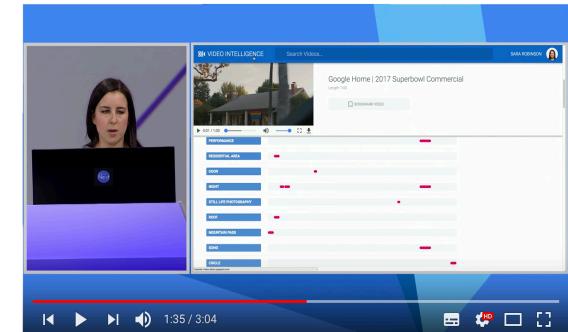


# Data Science



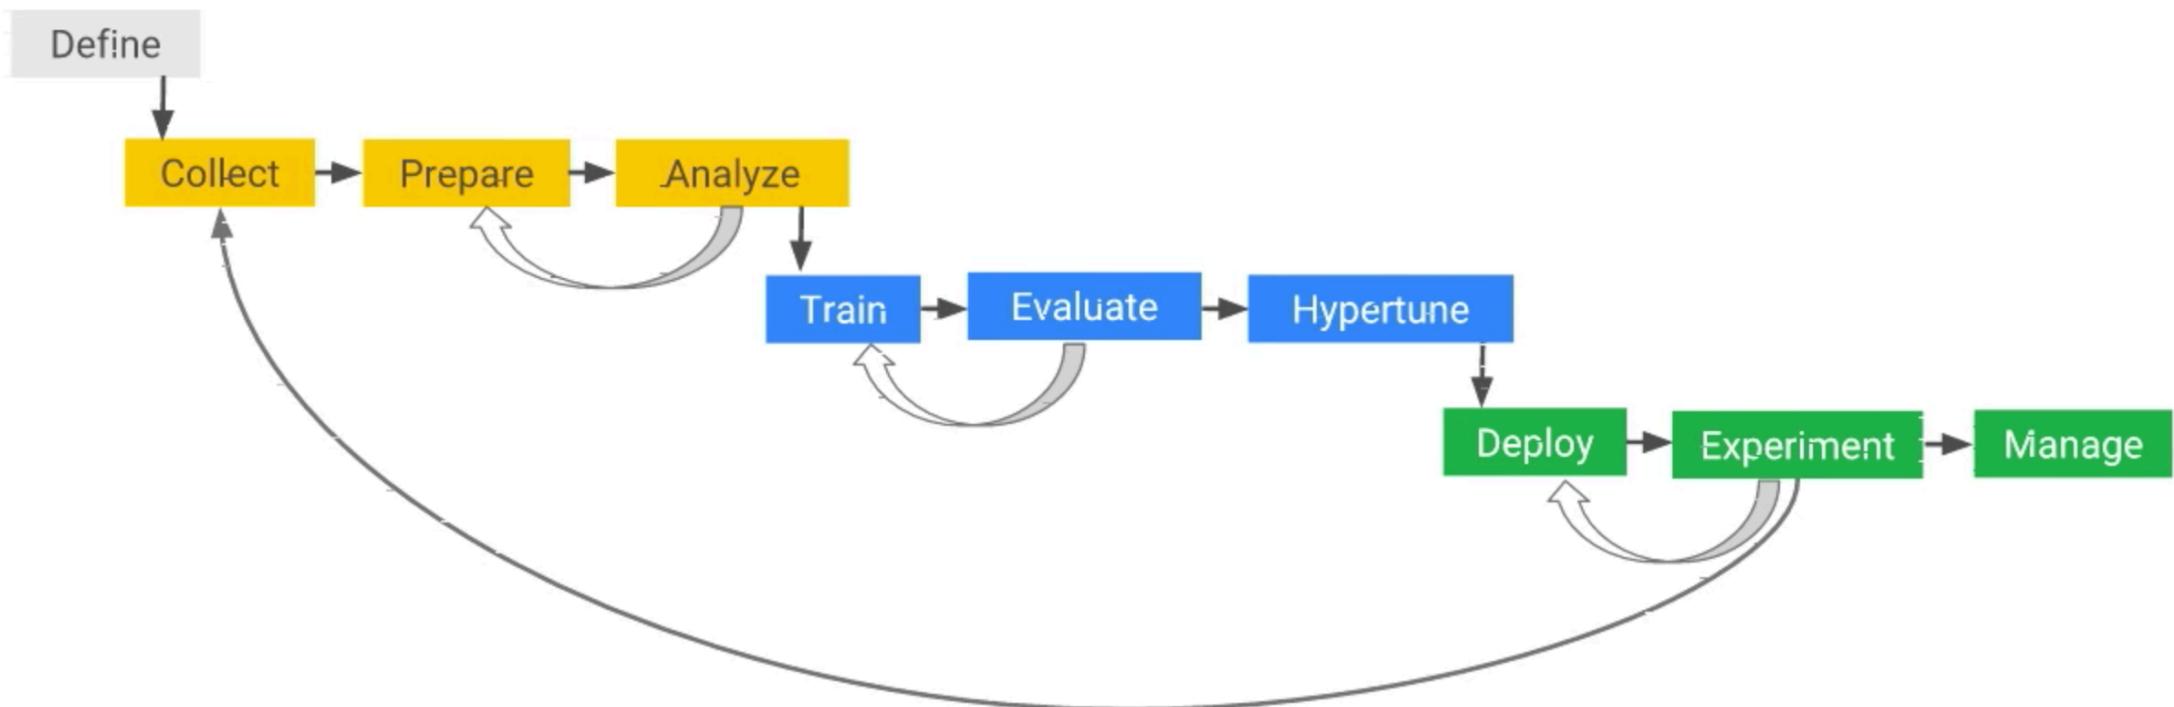
# Products

- Text Analysis
  - [uclassify.com](http://uclassify.com)
- Speech <-> Text Recognition
  - [Google Cloud](https://cloud.google.com/speech-to-text/)
  - [IBM Watson](https://www.ibm.com/watson/)
- Image and Video Classification
  - [clarifai.com](https://clarifai.com)
- Etc, etc, etc.



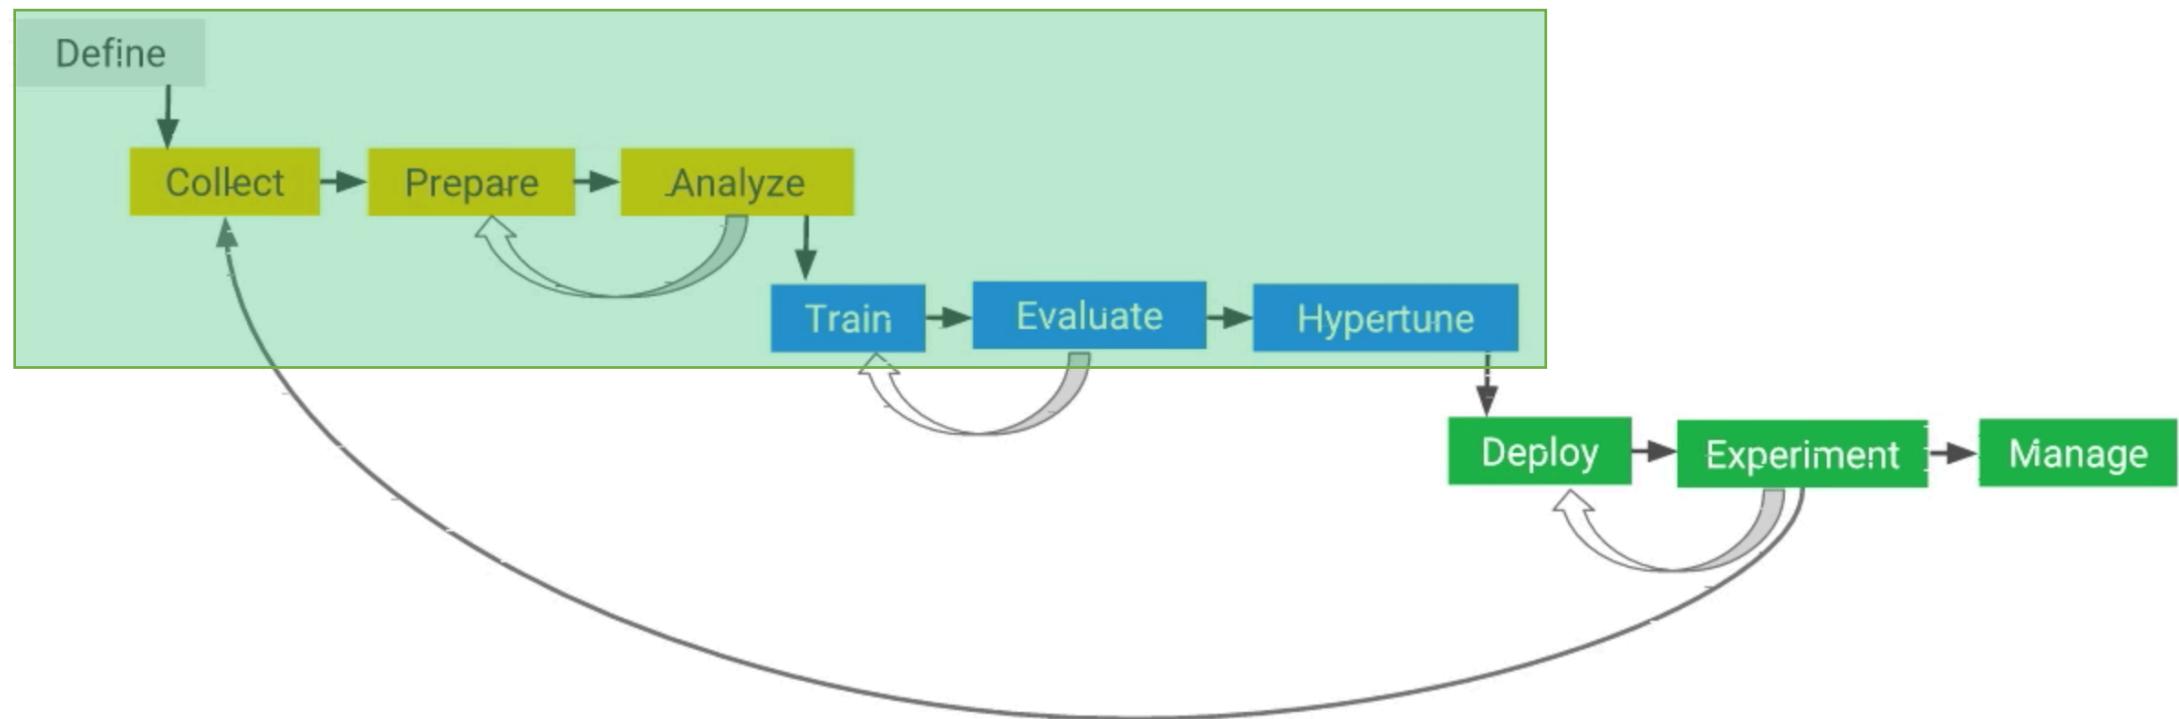
# ML Project Lifecycle

# ML Project Lifecycle

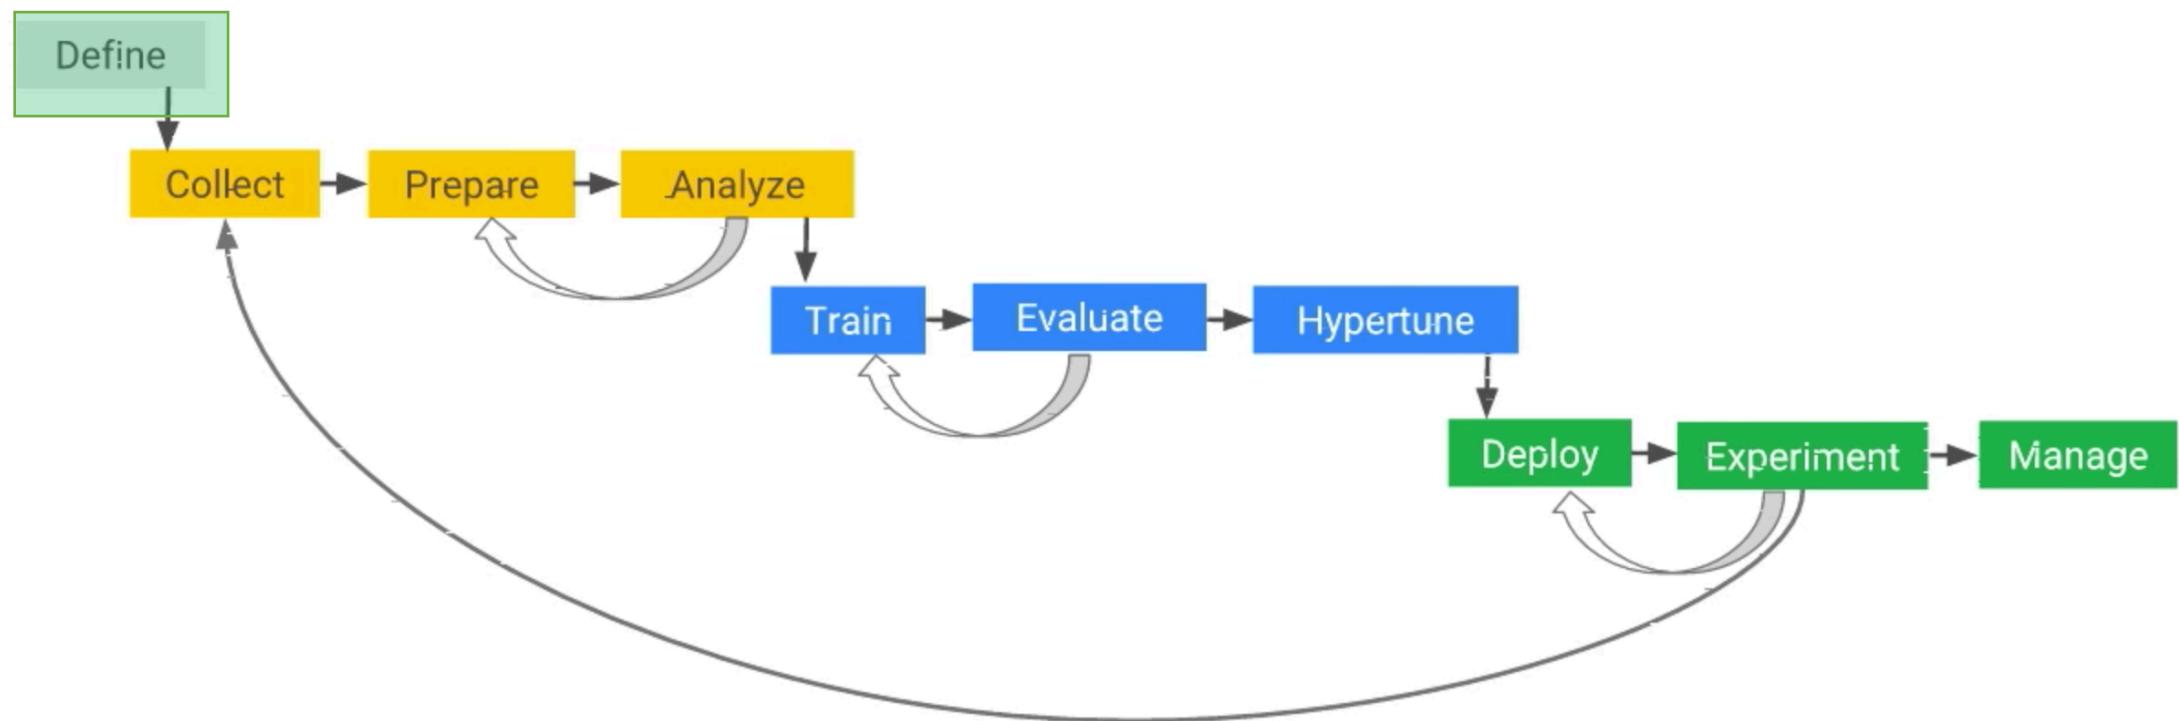


Google Cloud Next '17: <https://youtu.be/gkmAnu8DtIM?list=PLlivdWyY5sql8RuUibiH8sMb1Exlw0IAR>

# ML Project Lifecycle



# ML Project Lifecycle



# 0. Define

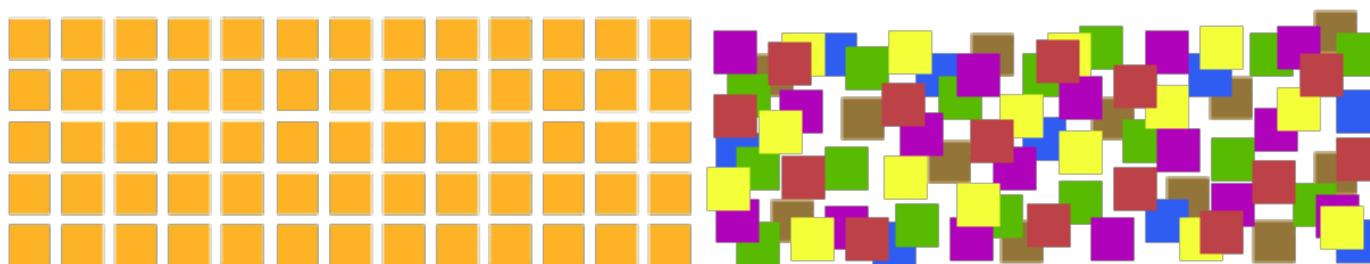
## • Data

- Rows/Example/Instance/Input/  
/Observation/Record/Point  
/Sample/Entity :  $x^{(i)}$
- Columns/Feature/Variable/Predictor  
/Characteristic/Field/Attribute :  $x_j^{(i)}$ 
  - Quantitative (numeric, continue)
  - Qualitative (textual, category)
- Dimension, Visualization
  - m Examples:  $i = 1..m$
  - n Features:  $j = 1..n$
- Label/class/output :  $y_i$ 
  - For each example (0/1)

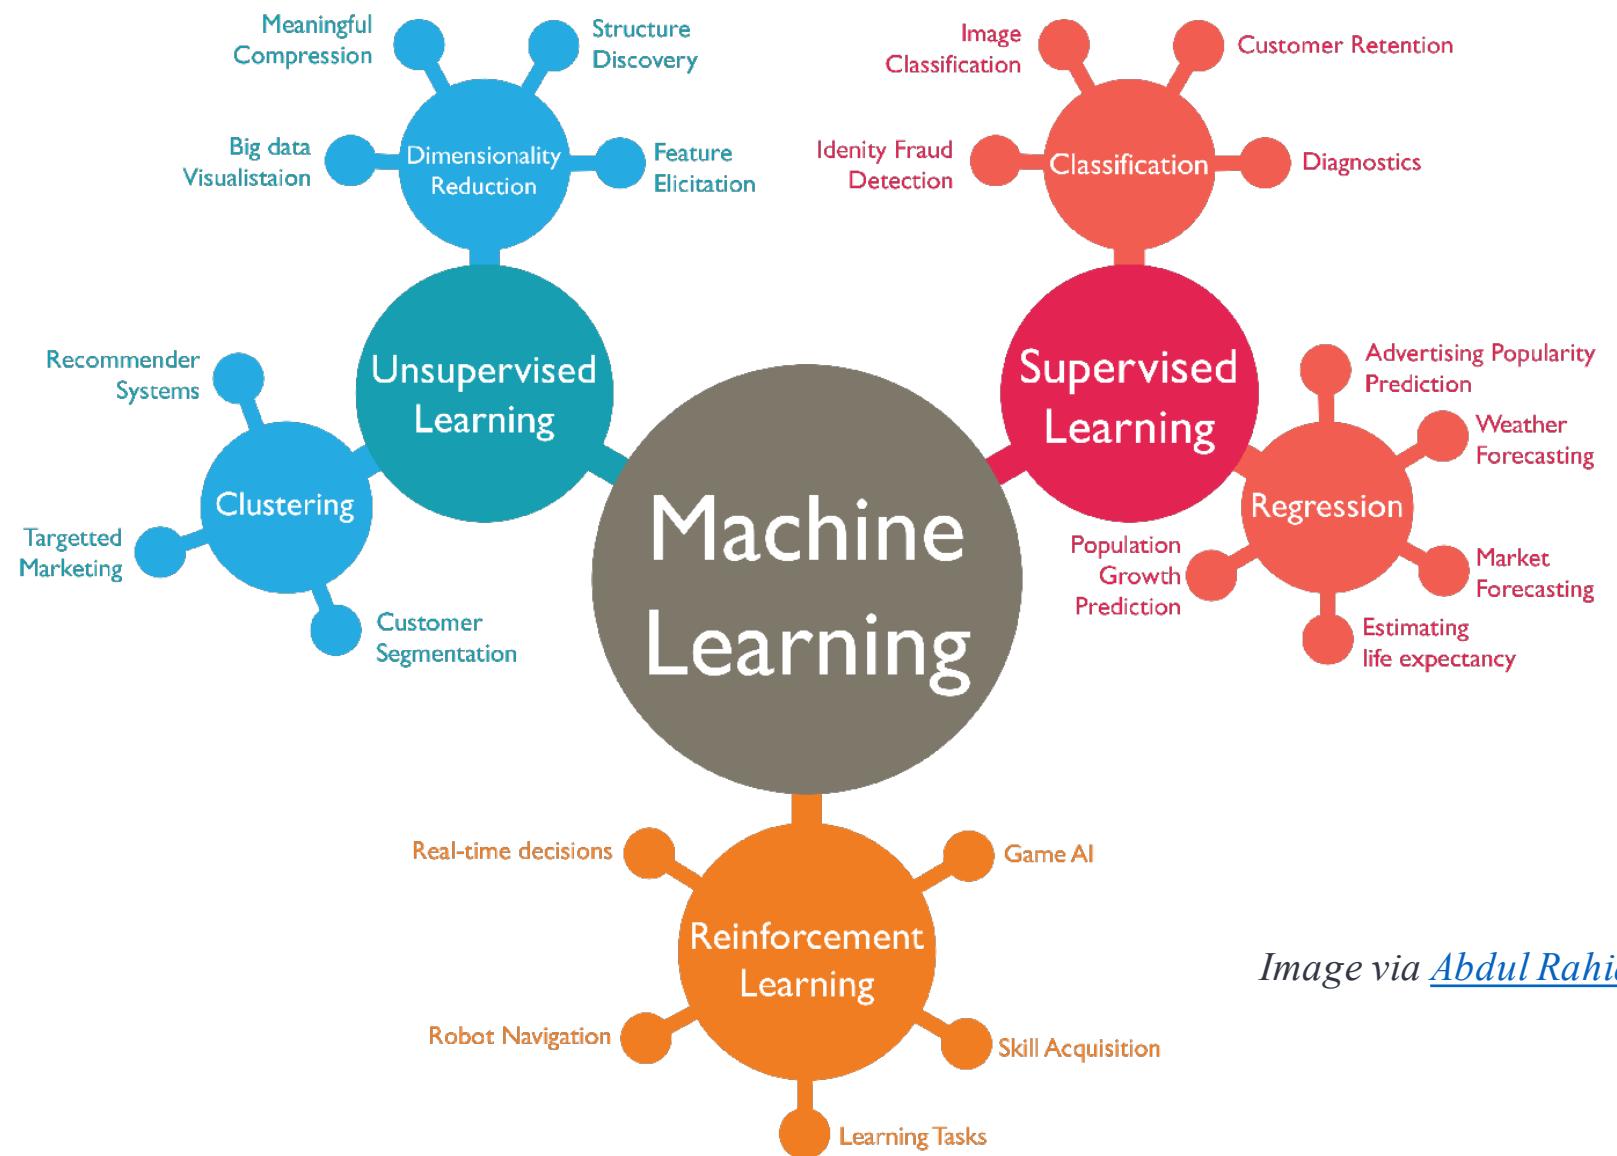
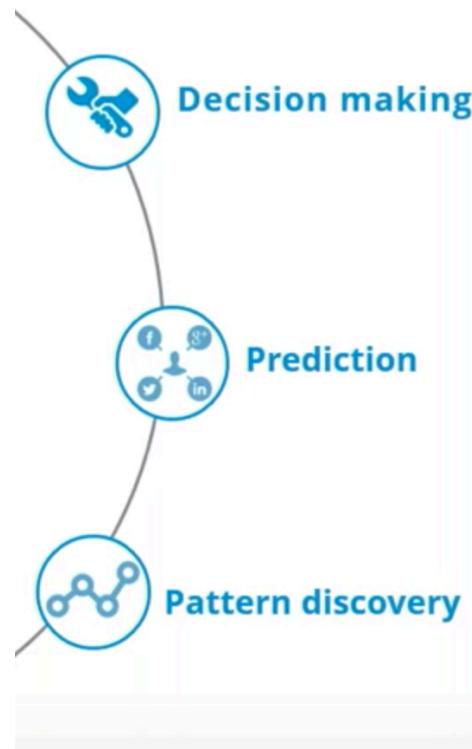
University	F1	F2	F3	F4	F5	F6	Region	Budget (MD)	Insertion (%)	Shanghai ranking?
U1	979	561	186	786	835	536	N	7	72.00	1
U2	895	247	985	206	870	246	N	7	67.00	1
U3	344	889	643	951	783	162	E	2	35.00	0
U4	400	959	999	312	981	254	W	8	57.00	0
U5	243	521	393	596	400	138	E	6	29.00	1
U6	882	722	518	541	425	551	W	10	80.00	1
U7	814	193	829	192	334	597	E	8	68.00	0
U8	186	972	763	968	217	772	S	1	20.00	0
U9	647	656	527	393	738	813	S	7	65.00	1
U10	568	570	458	799	682	530	N	4	59.00	1

# 0. Define

- **Data**
  - Structured
    - CSV, XML, JSON, XLSX, etc.
  - Unstructured
    - DOC, HTML, PDF, PNG, MP3, MP4, etc.

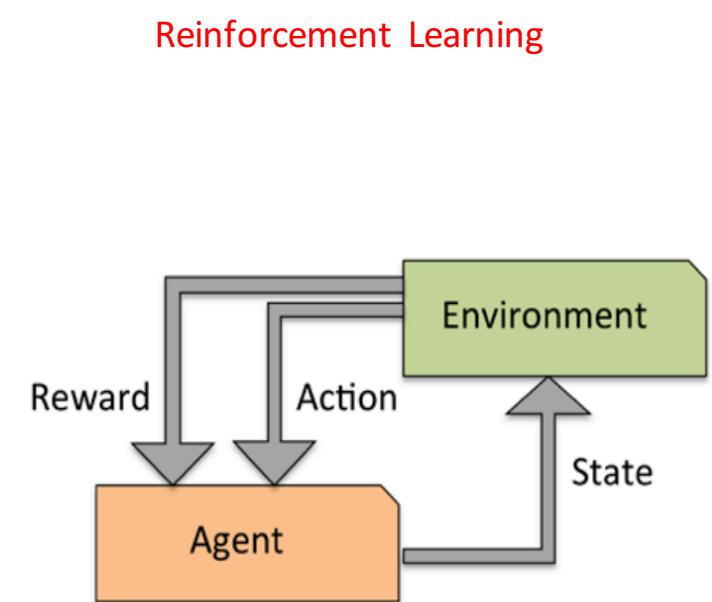
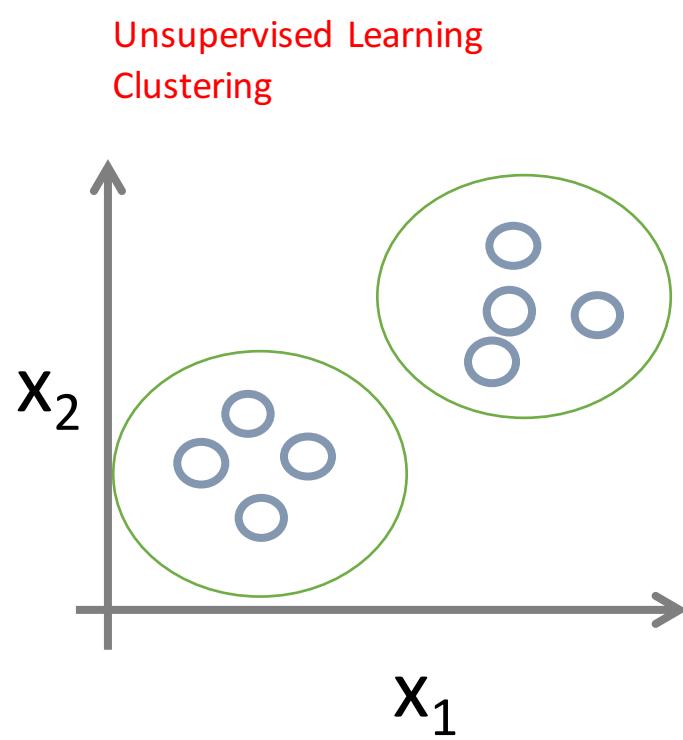
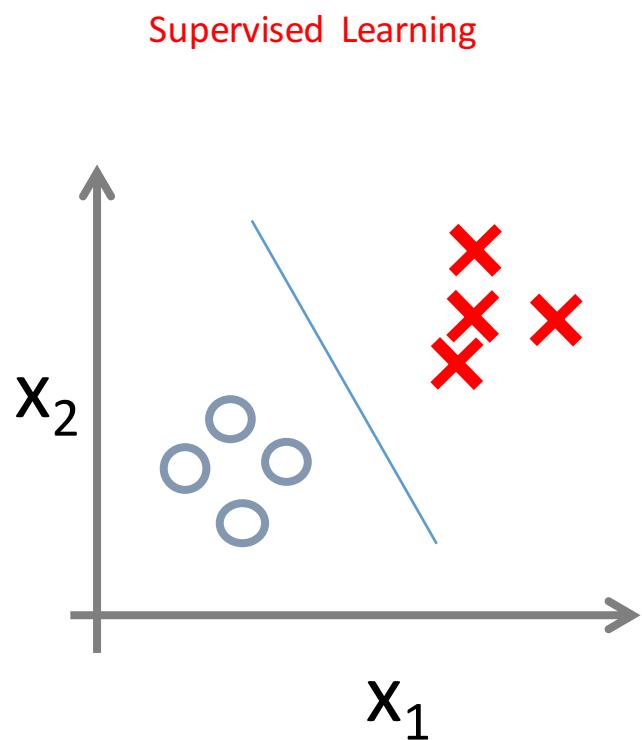


# 0. Define

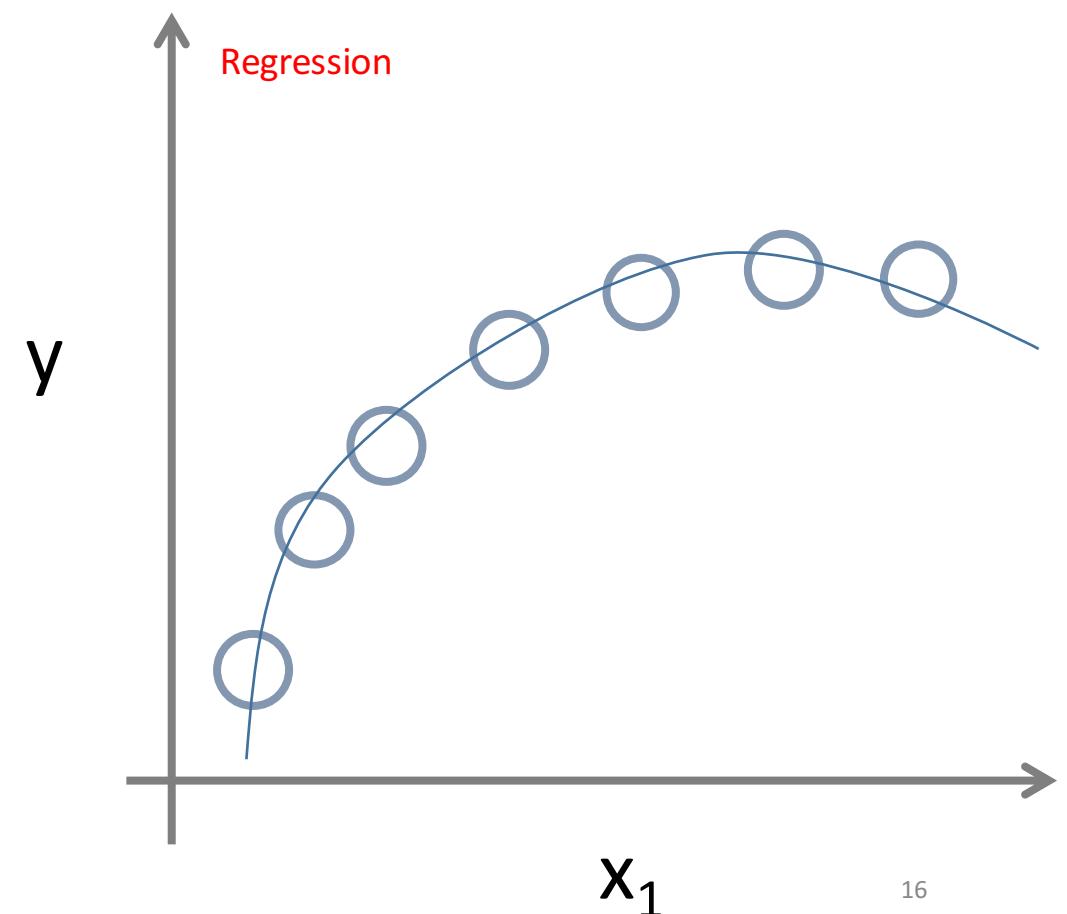
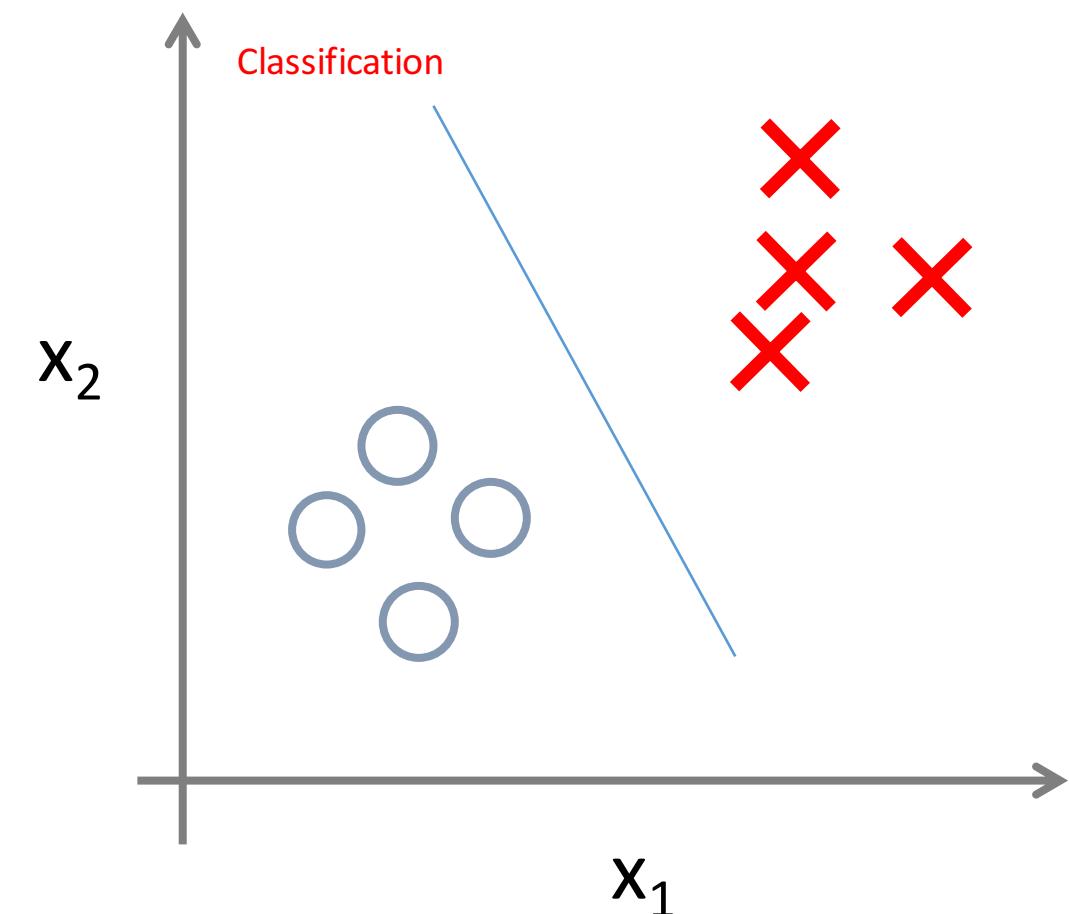


*Image via [Abdul Rahid](#)*

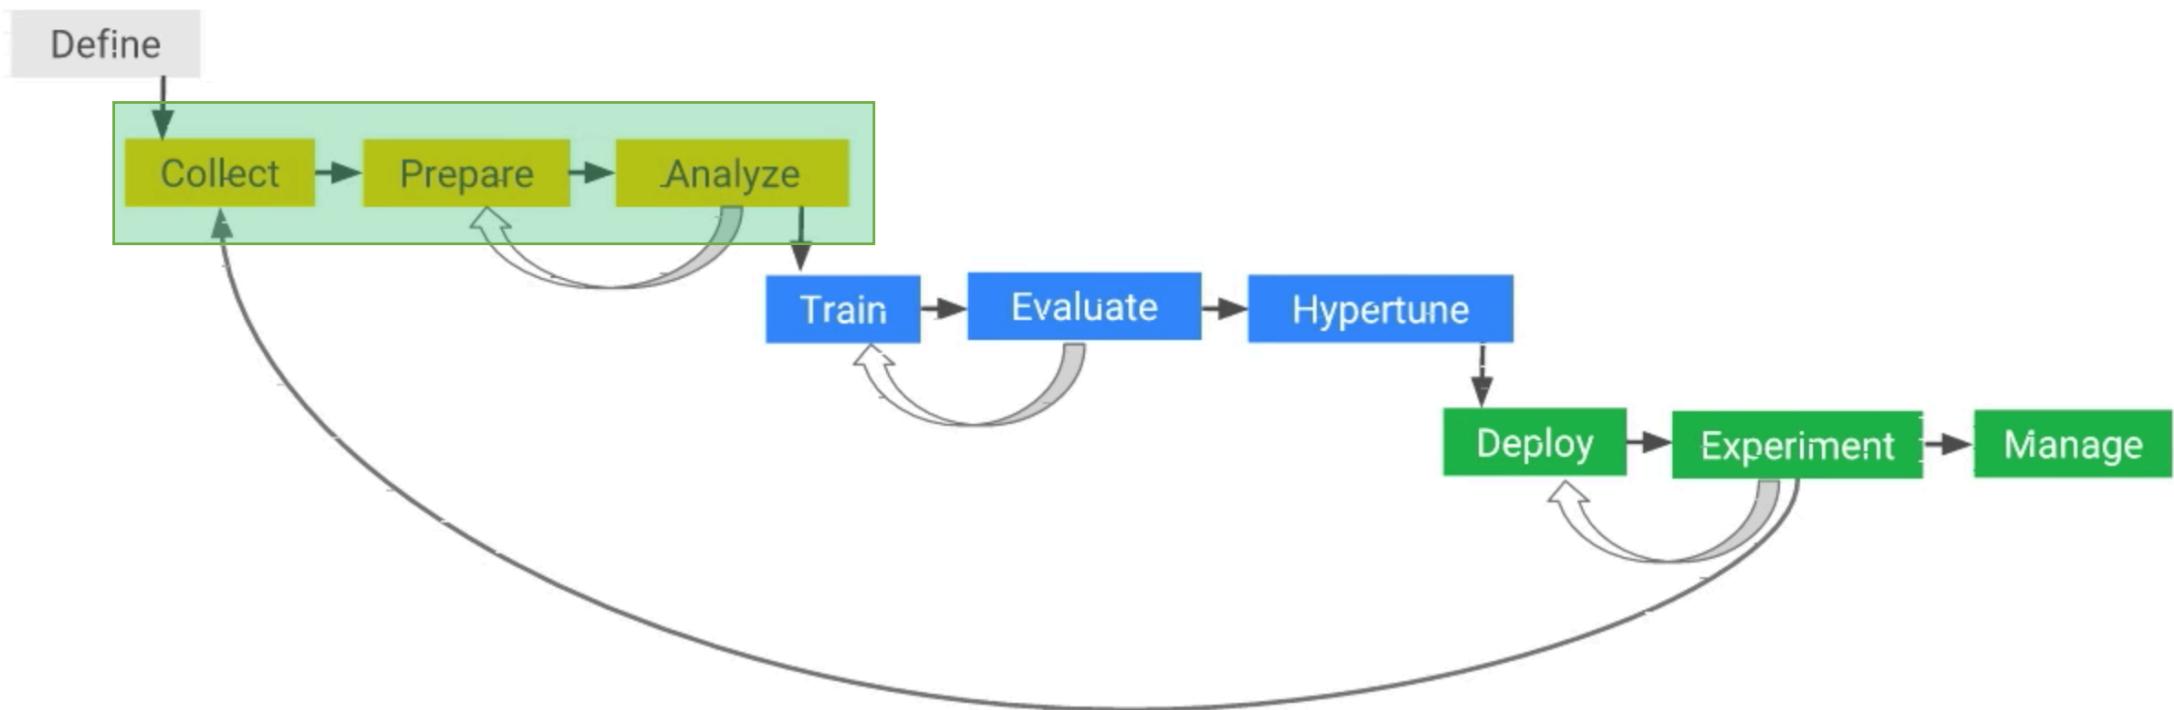
# 0. Define



## 0. Define



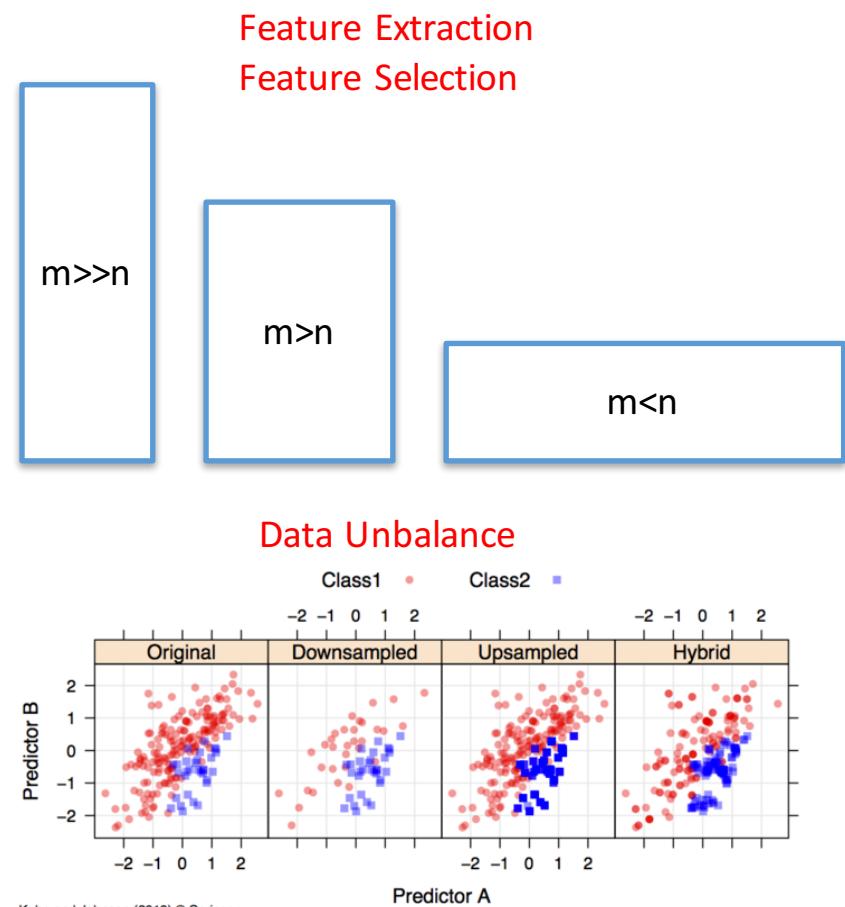
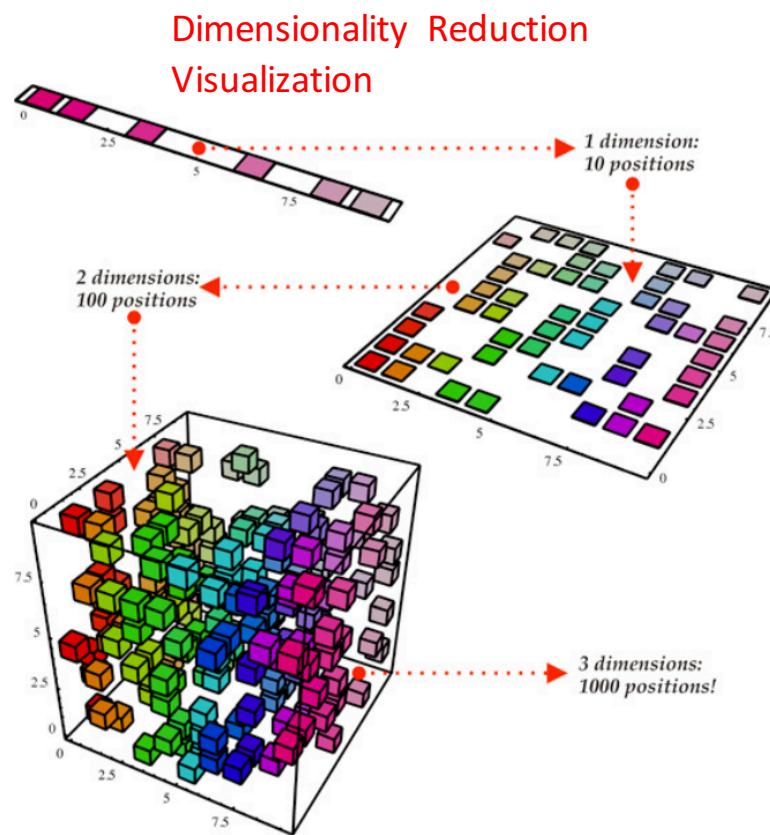
# ML Project Lifecycle



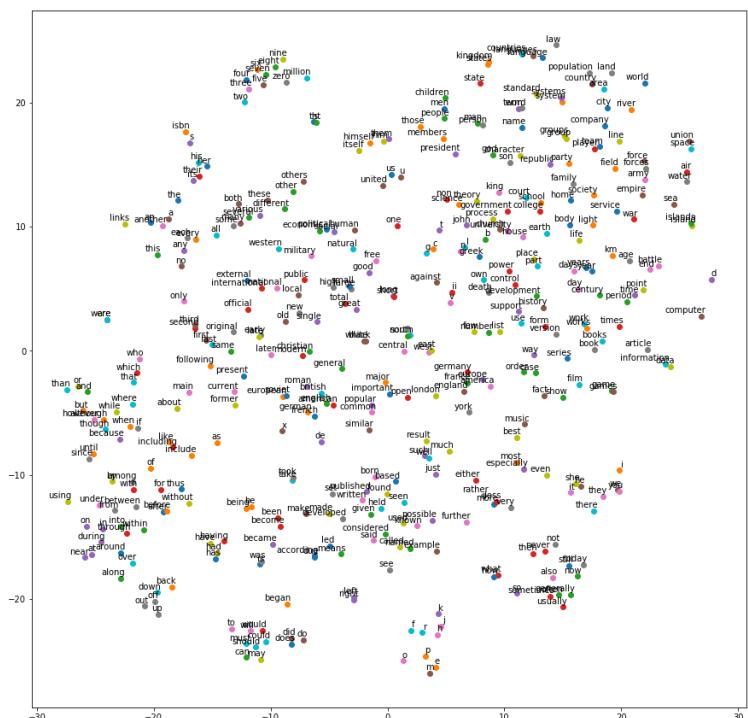
# 1. Collect

- Government, NGOs, Scientific Communities, etc.
- Open Data Platforms,
- Web data: AWS, Twitter API, Facebook API, Scraping tools, etc.
- Mobile Data: google maps, logs, etc
- Enterprise data: Databases, Datawarehouses
- IoT, Industrial IoT data
- More: Forbes: [Big Data: 33 Brilliant And Free Data Sources Anyone Can Use](#)

## 2. Prepare

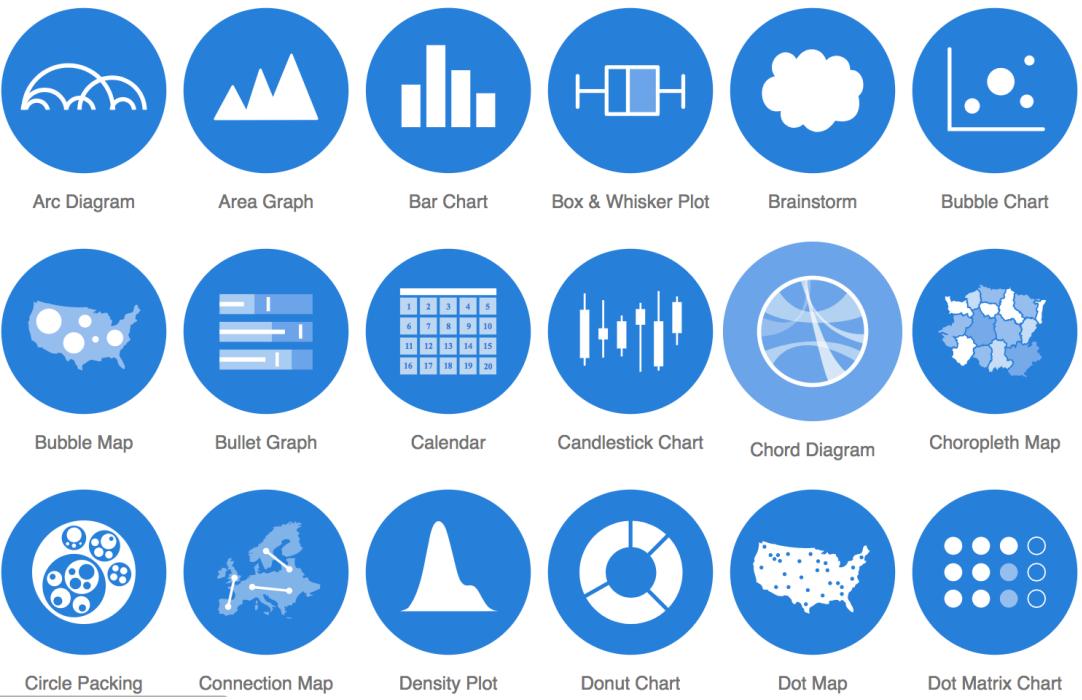


### 3. Analyze



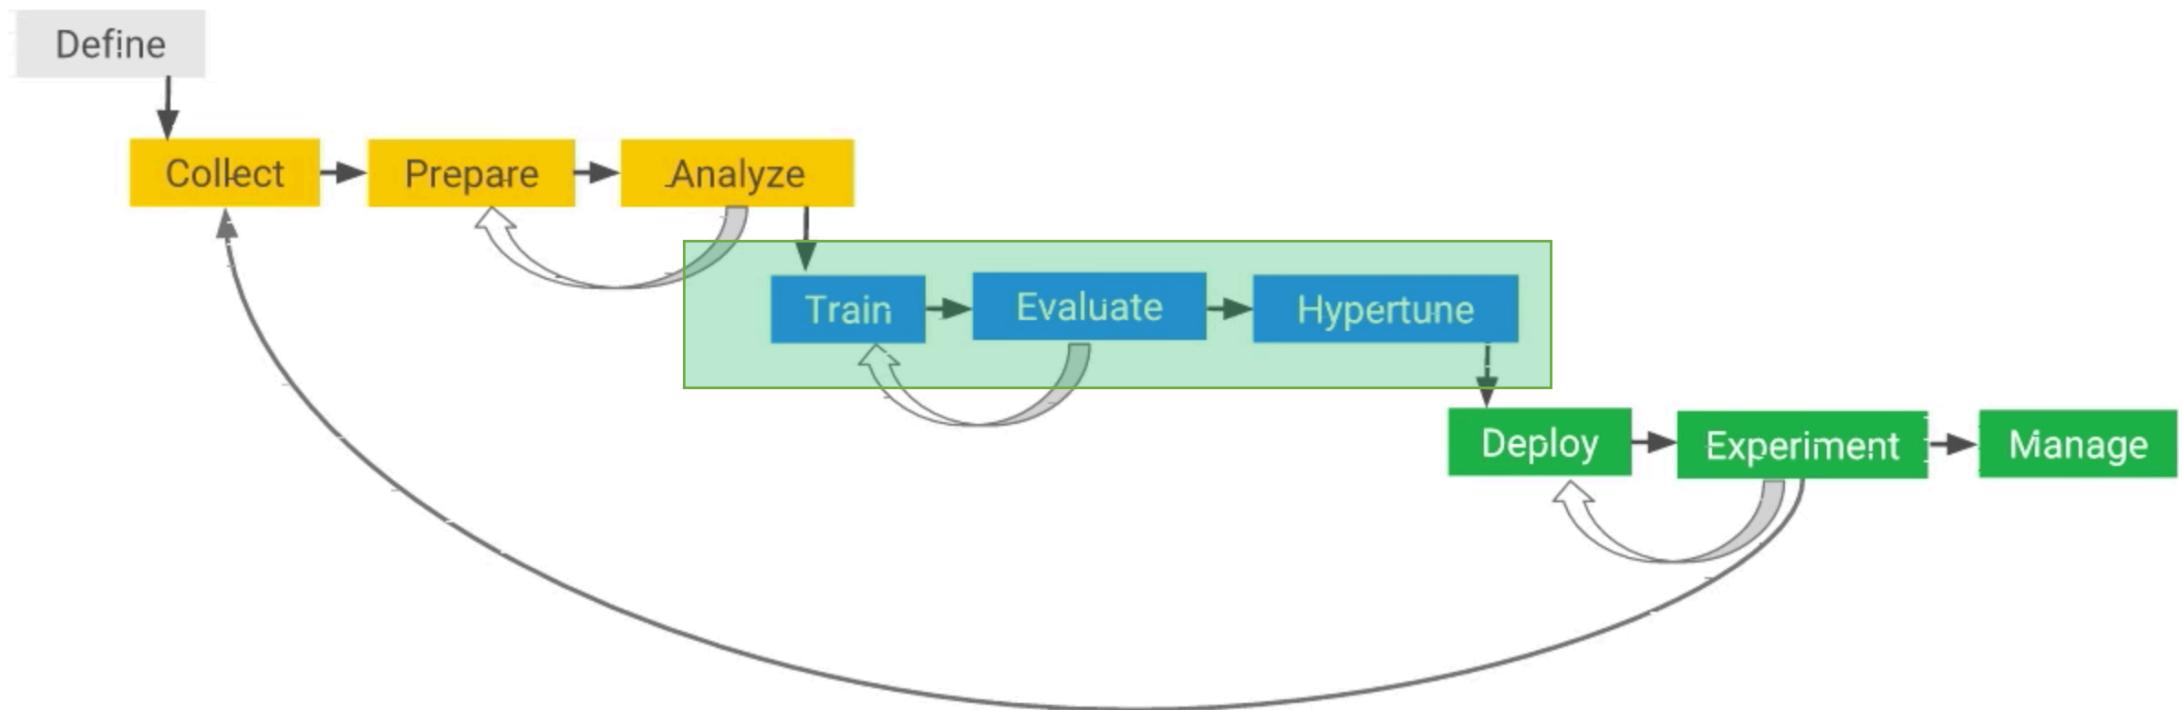
## T-SNE

## Stats , Transform, Visualize

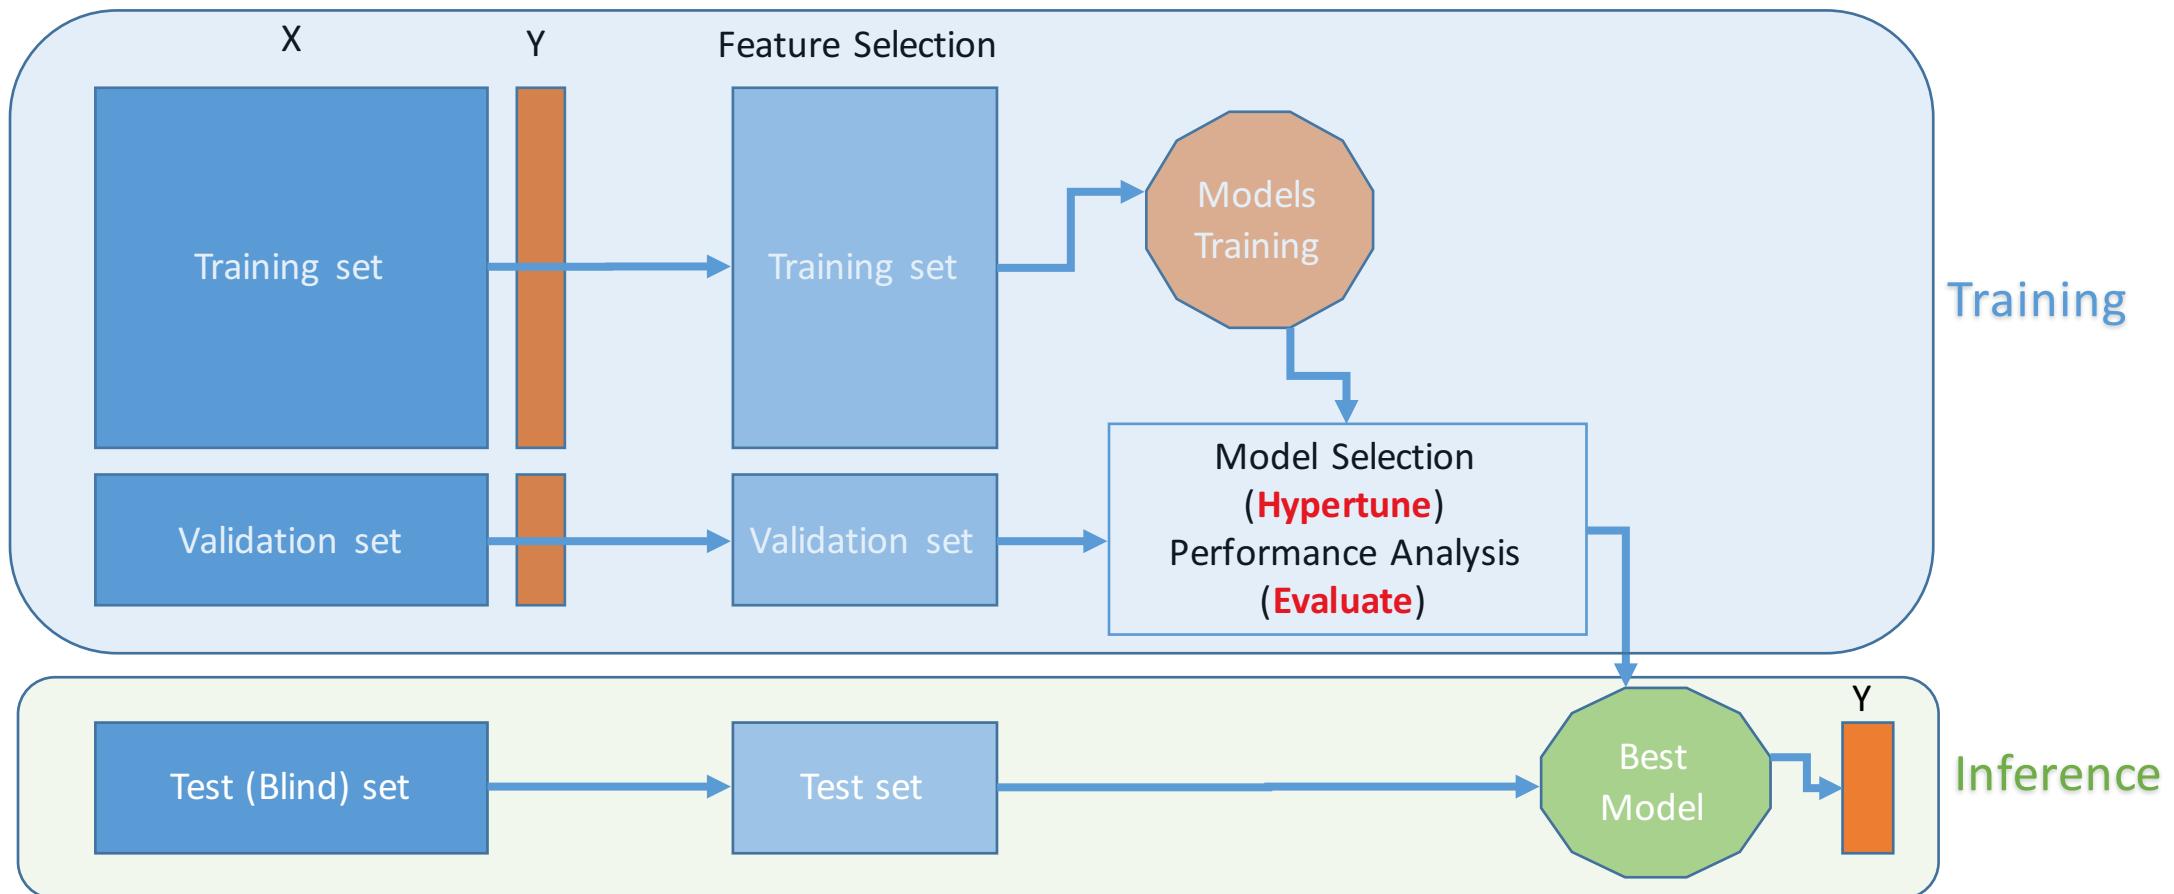


<https://datavizcatalogue.com/index.html>

# ML Project Lifecycle



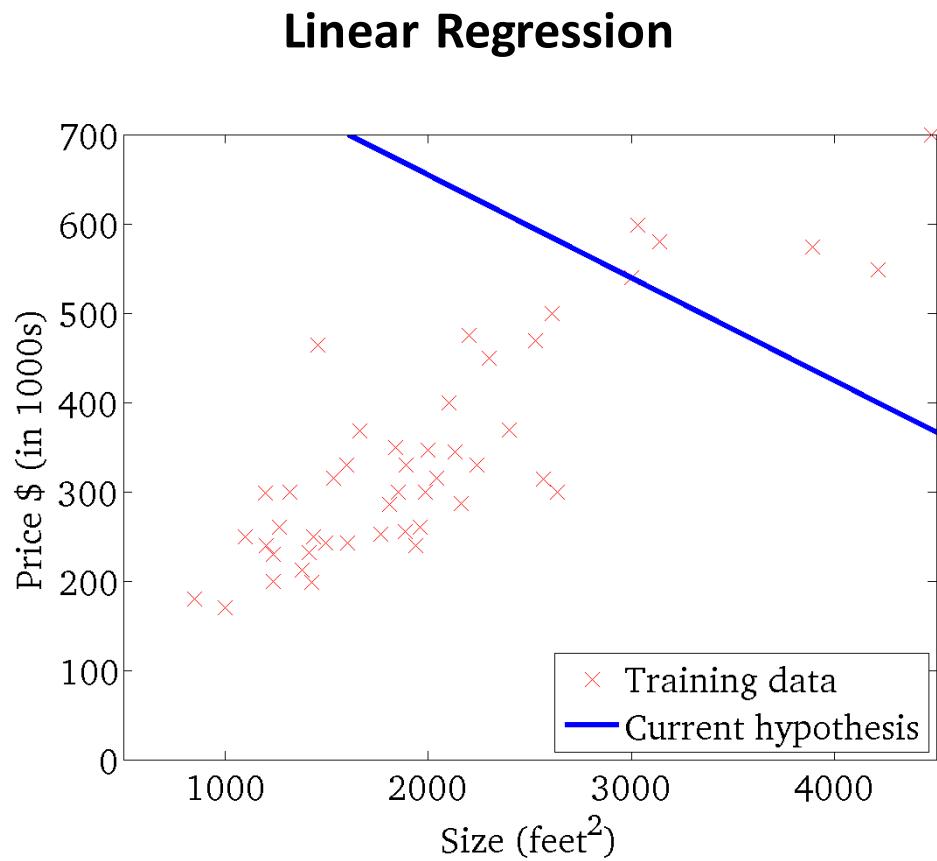
## 4. Train



# 4. Train



# 4. Train



Hypothesis

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$h_{\theta}(x) = \theta^T x$$

Cost function

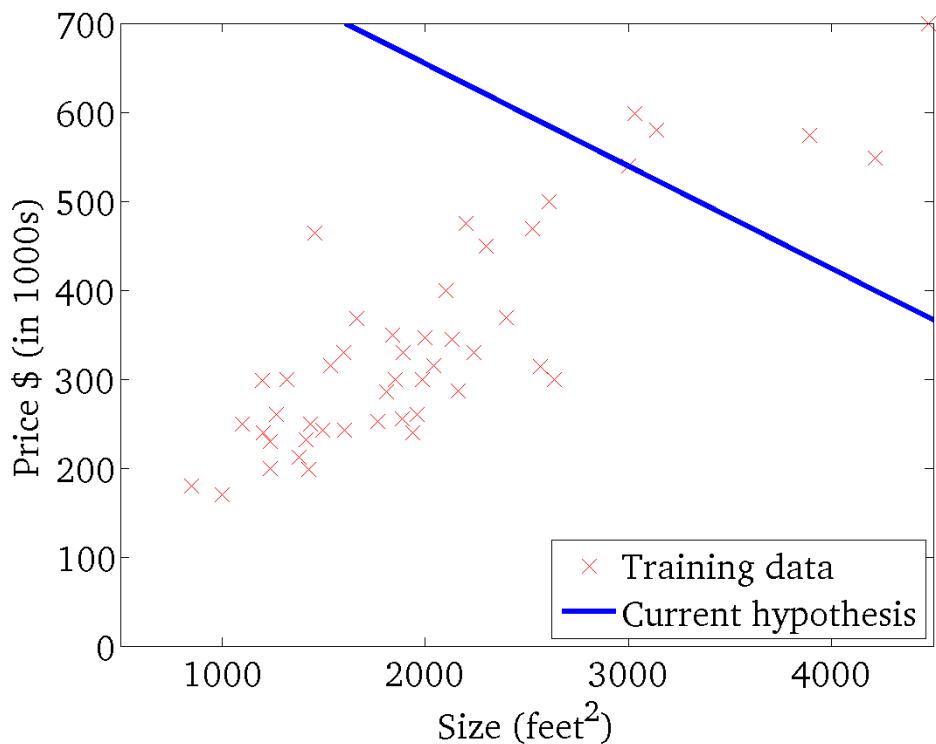
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Optimization

$$\min_{\theta} J(\theta)$$

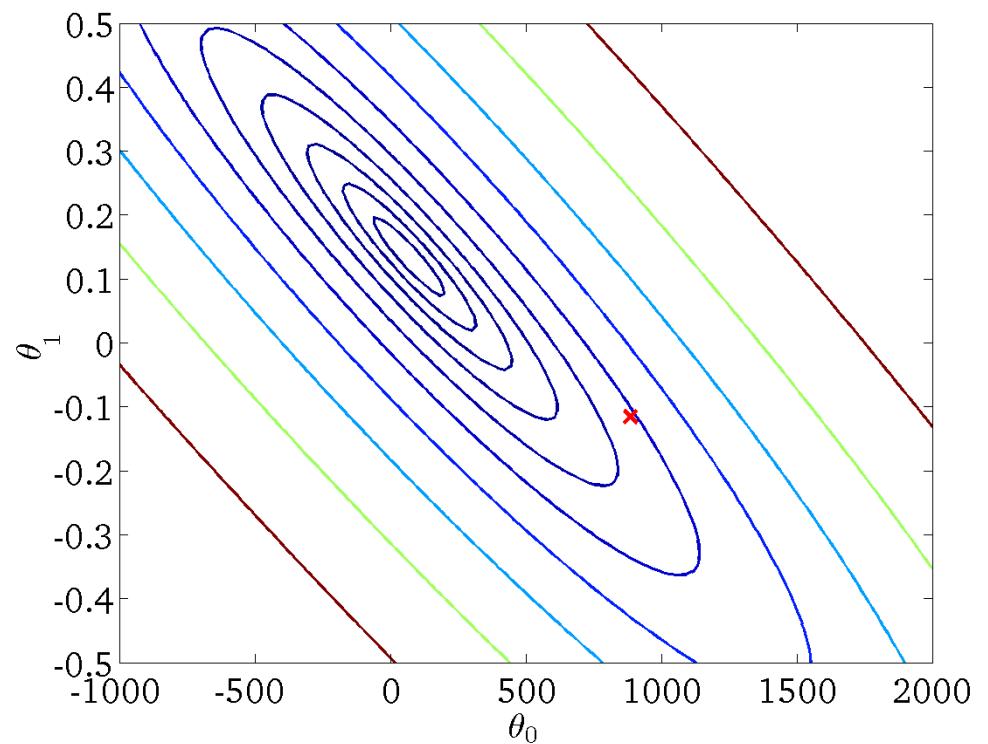
$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



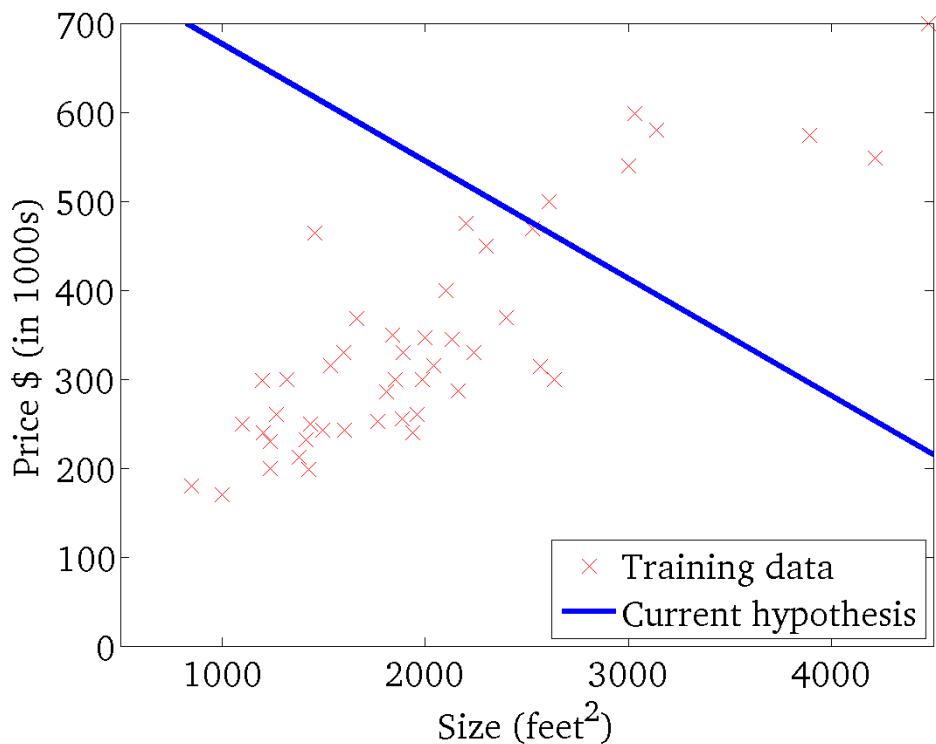
$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



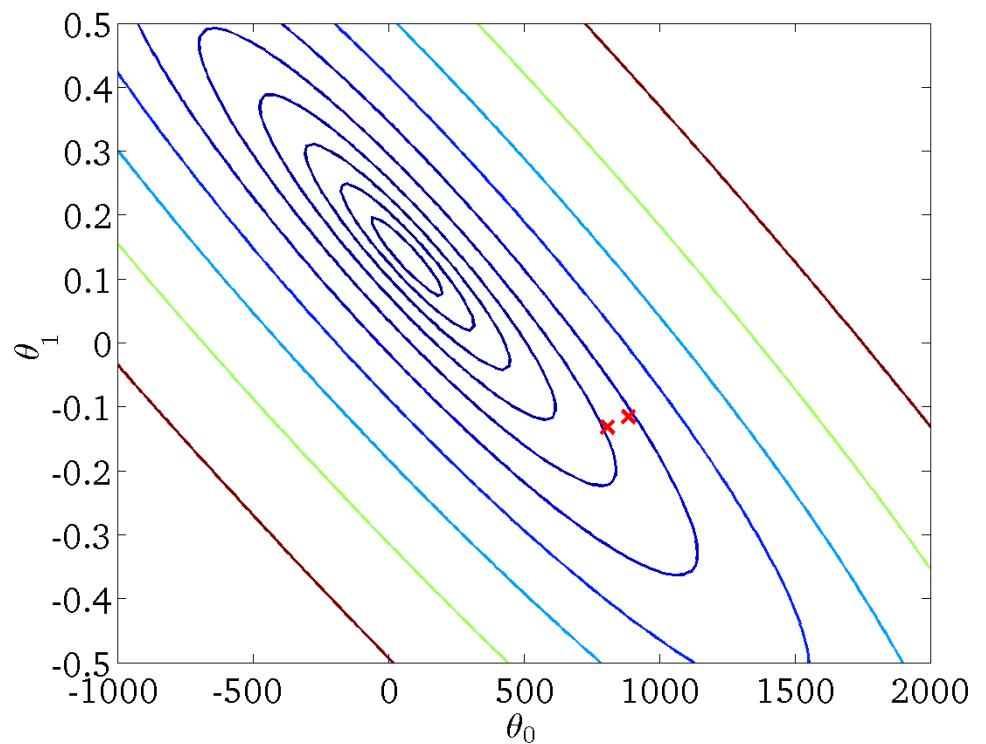
$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$  this is a function of  $x$ )



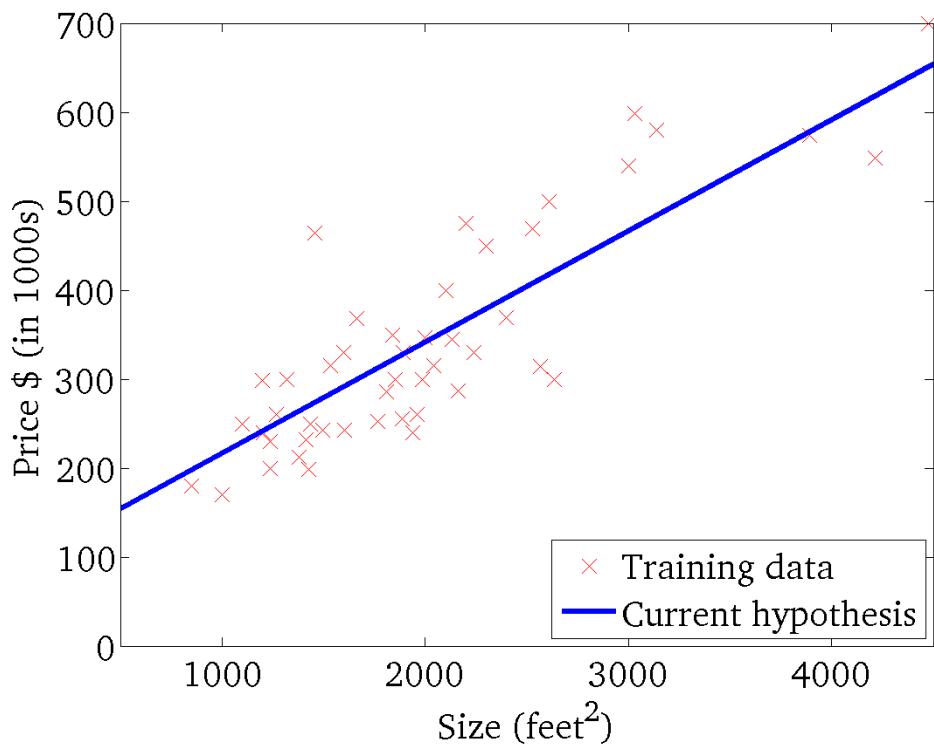
$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



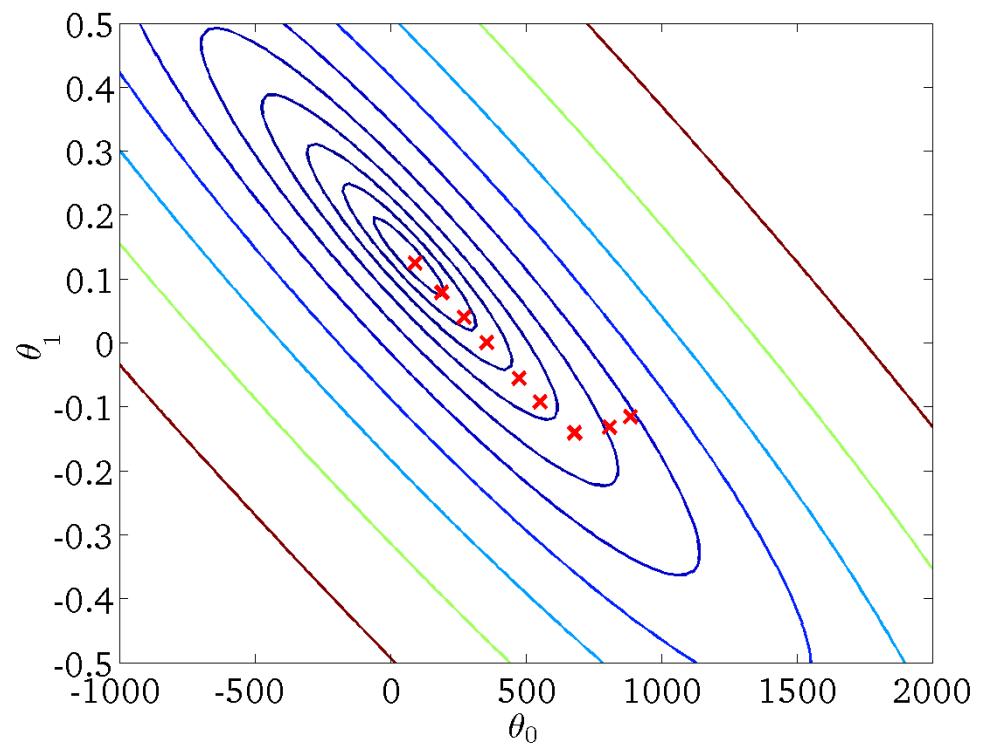
$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$  this is a function of  $x$ )



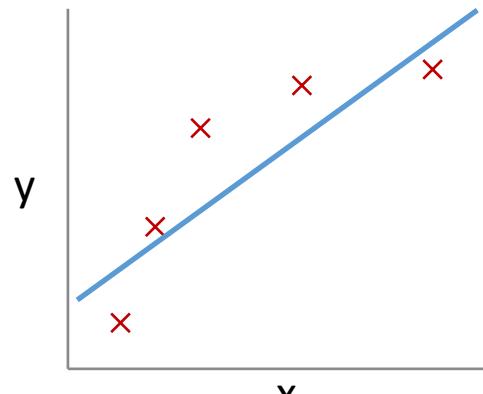
$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



# 4. Train

## Linear Regression



$$h_{\theta}(x) = \theta^T x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\min_{\theta} J(\theta)$$

Regularization  
term

## 4. Train

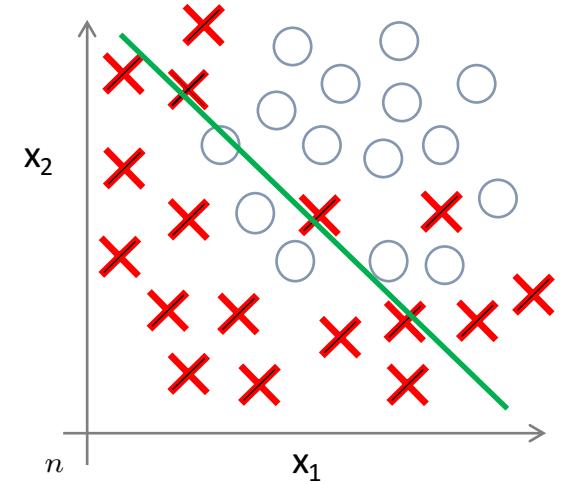
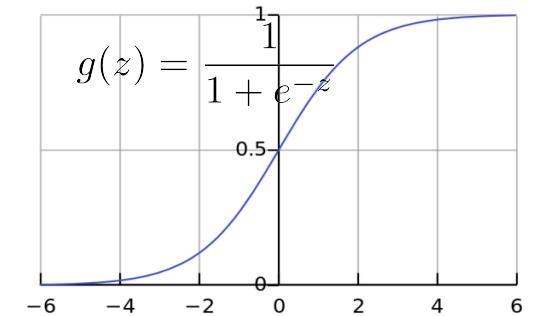
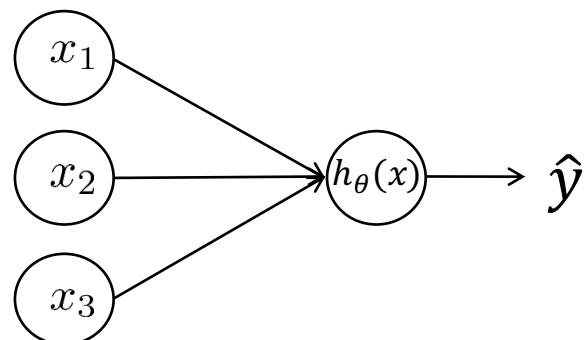
### Logistic Regression

Want  $0 \leq h_\theta(x) \leq 1$

$$h_\theta(x) = g(\theta^T x)$$

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\min_{\theta} J(\theta)$$



Regularization term

# 4. Train

**ANN**

$$z_1^{[1]} = w_1^{[1]T} x + b_1^{[1]}, \quad a_1^{[1]} = g(z_1^{[1]})$$

$$z_2^{[1]} = w_2^{[1]T} x + b_2^{[1]}, \quad a_2^{[1]} = g(z_2^{[1]})$$

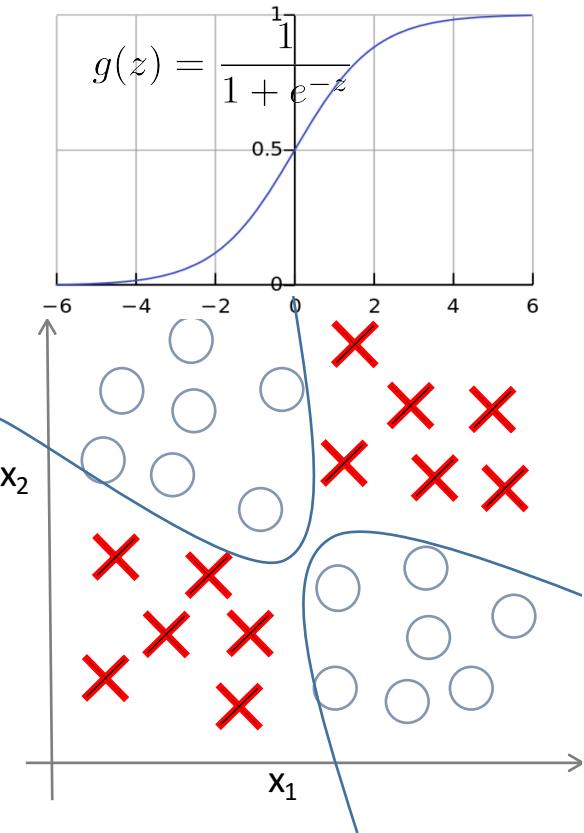
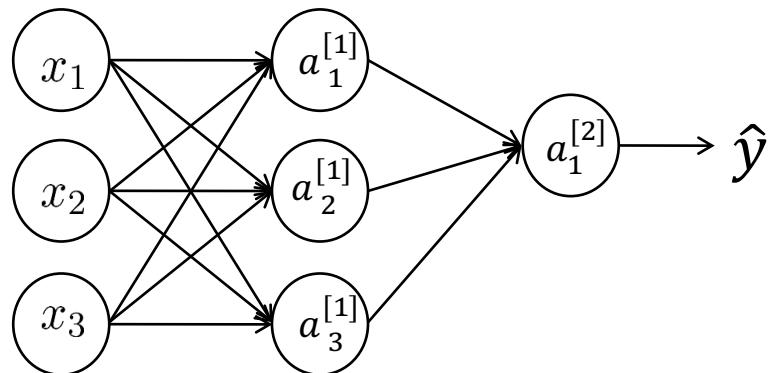
$$z_3^{[1]} = w_3^{[1]T} x + b_3^{[1]}, \quad a_3^{[1]} = g(z_3^{[1]})$$

$$z_1^{[2]} = w_1^{[2]T} x + b_1^{[2]}, \quad a_1^{[2]} = g(z_1^{[2]})$$

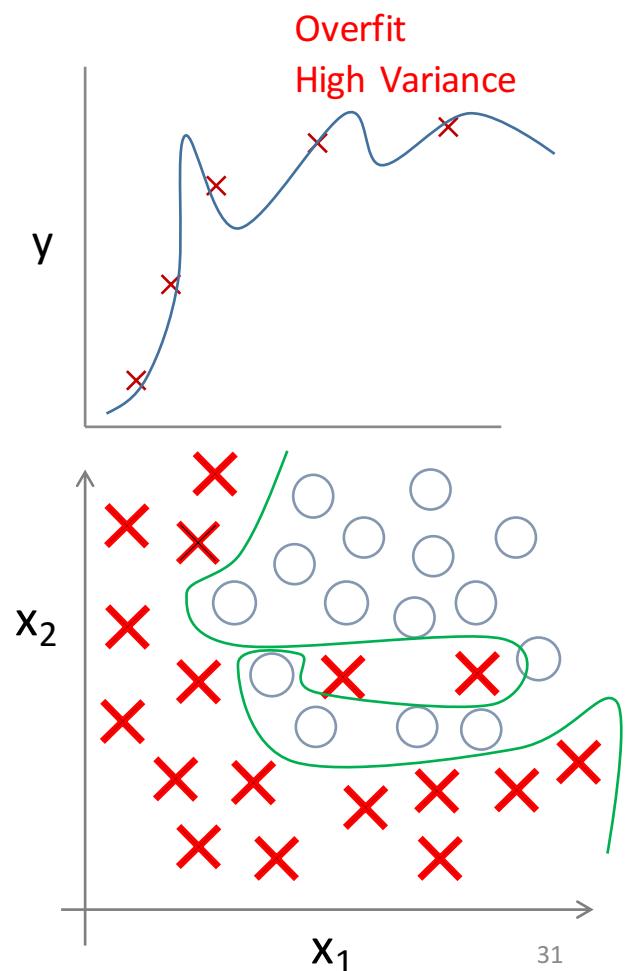
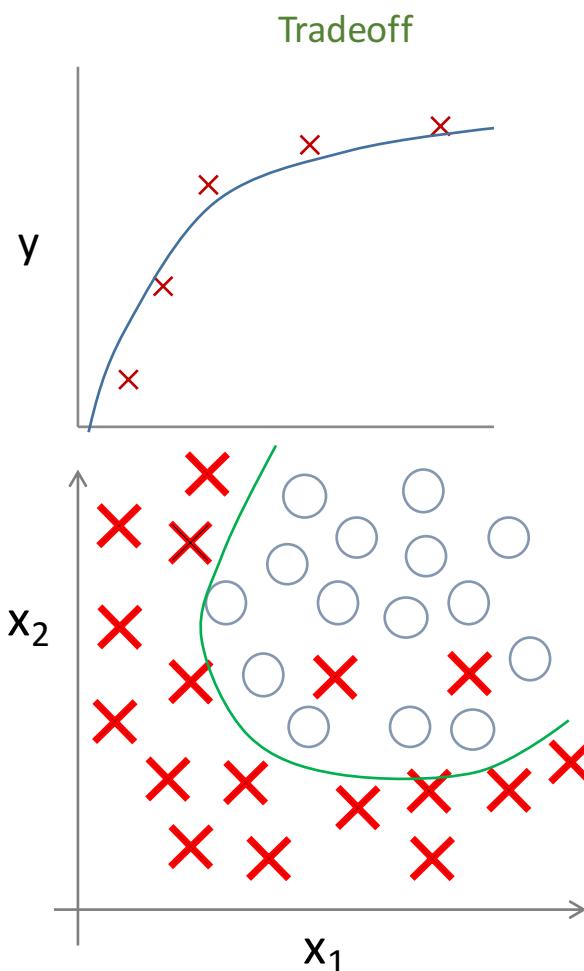
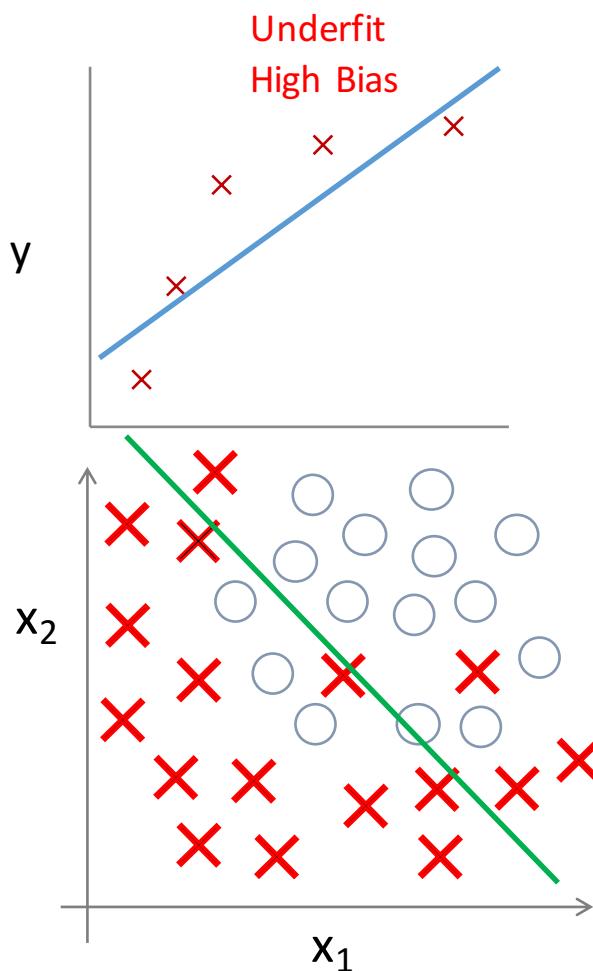
$$J = -\frac{1}{m} \sum_{i=0}^m \left( y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)}) \right) + \frac{\lambda}{2m} \sum_{j=1}^n \sum_{l=1}^L W_j^{l2}$$

Regularization term

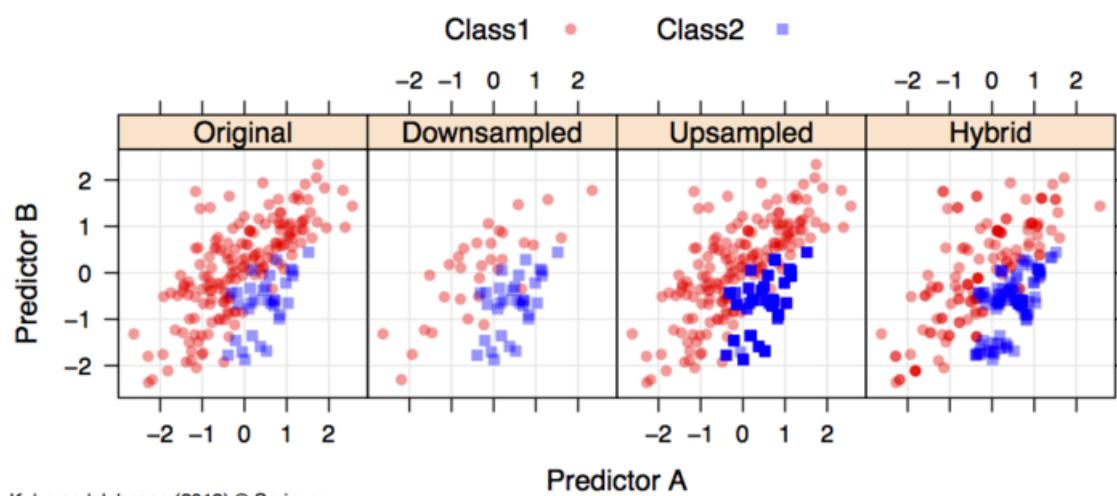
$$\min_{W,b} J(W, b)$$



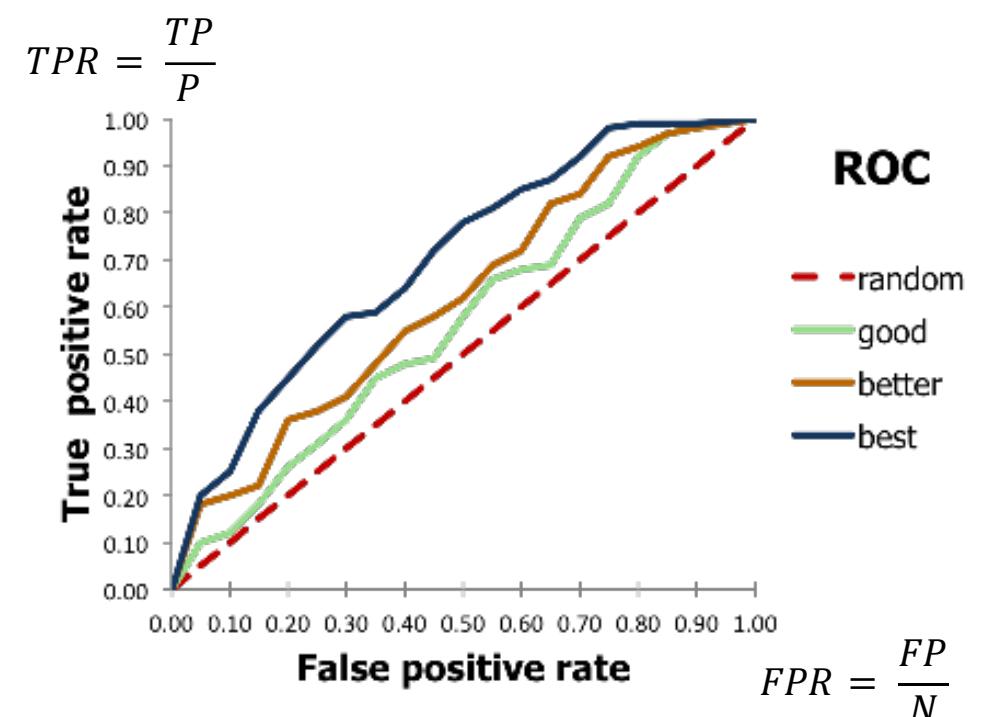
## 5. Evaluate



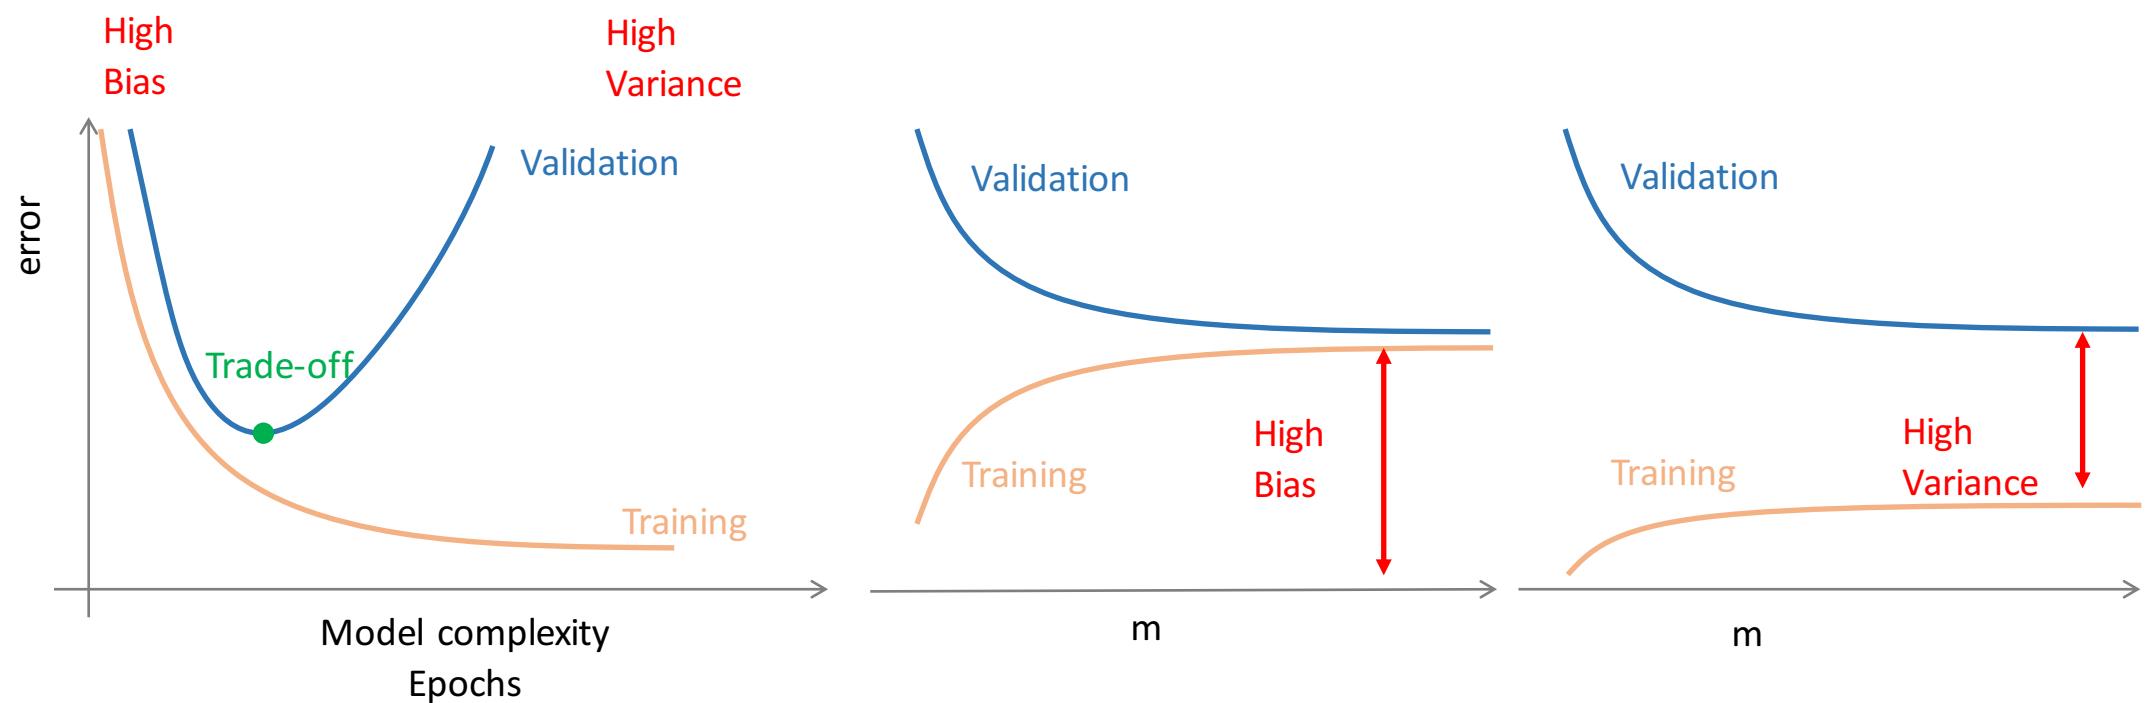
# 5. Evaluate



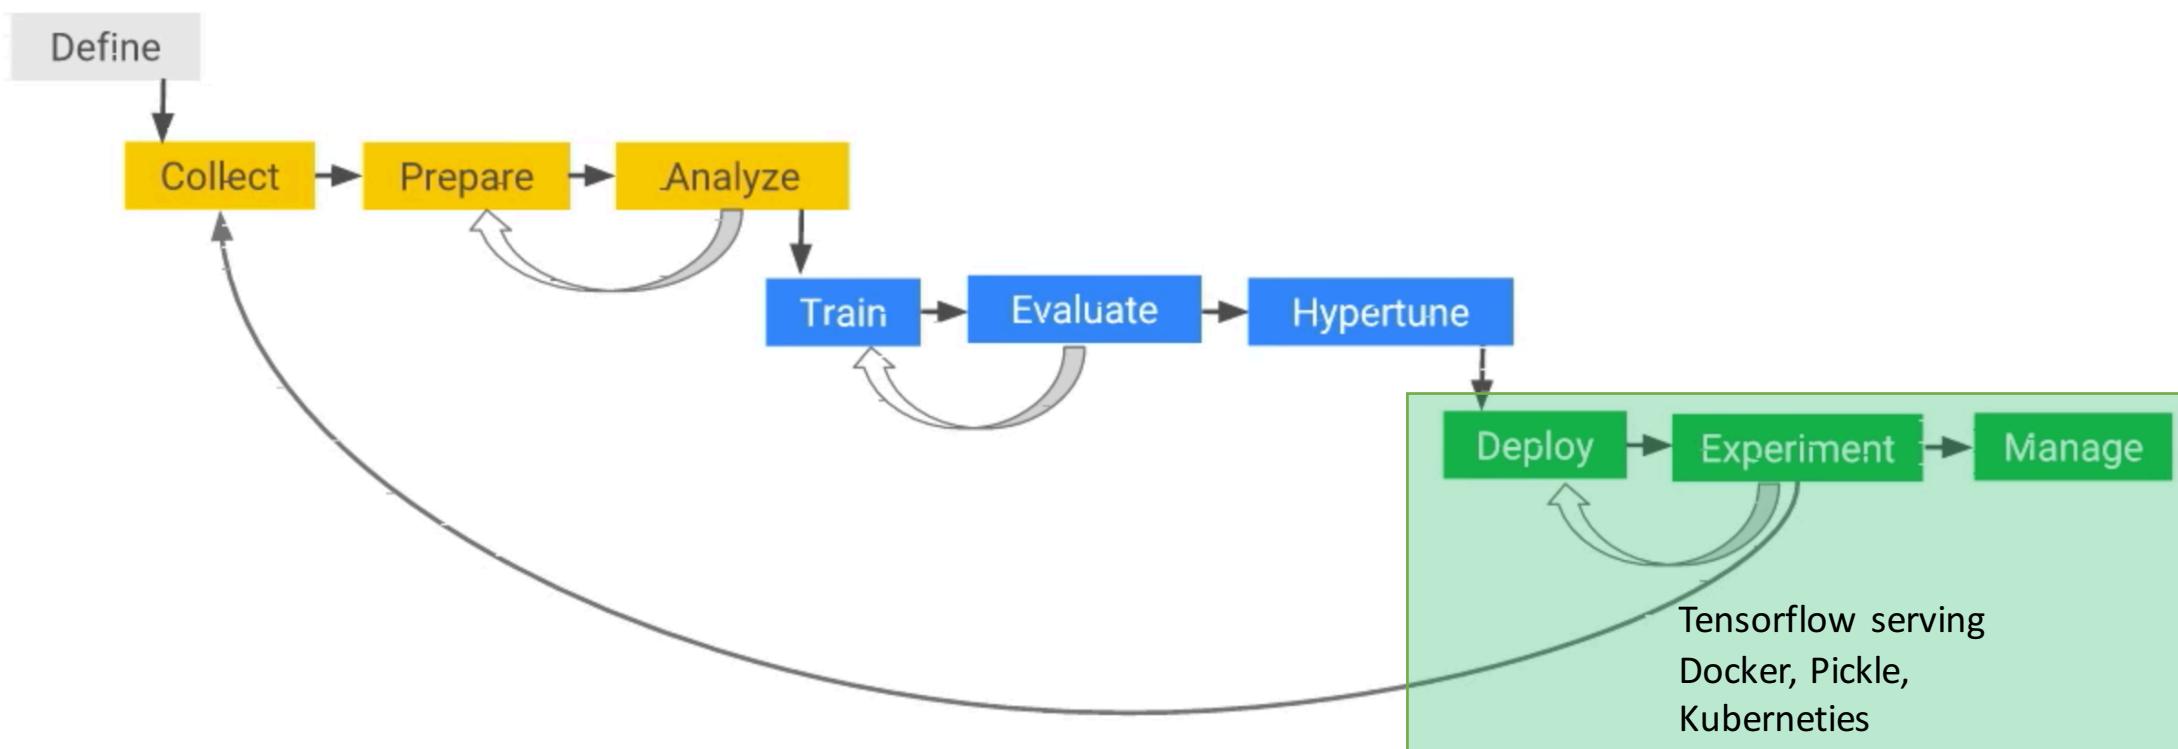
Kuhn and Johnson (2013) © Springer



# 7. Hypertune

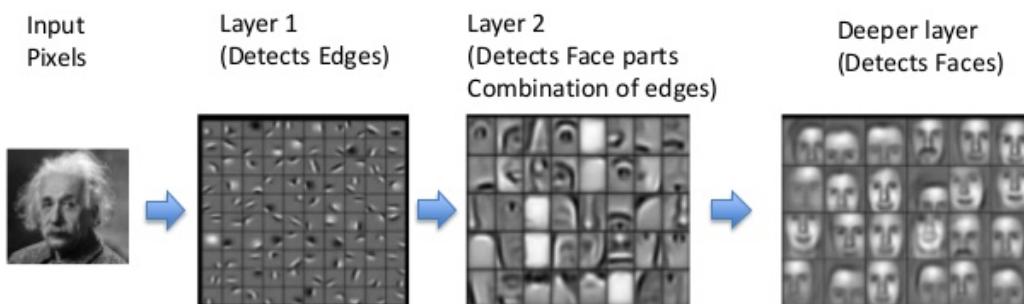
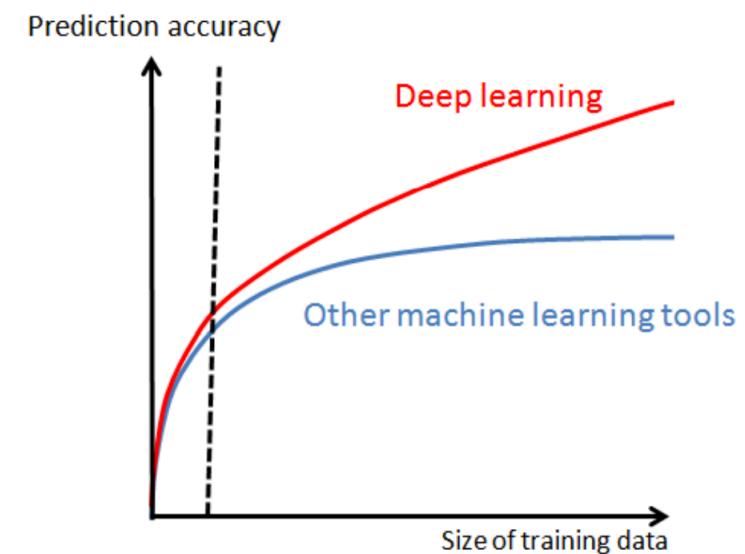
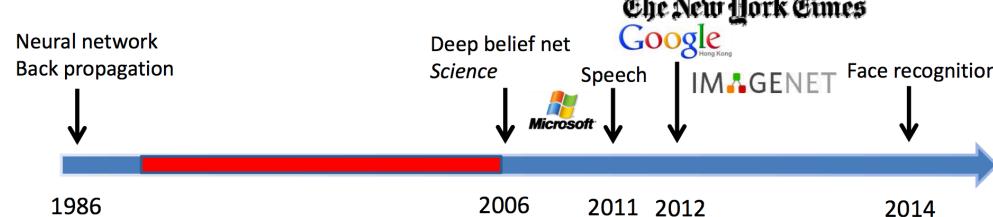


## 8. Deploy, 9. Experiment, 10. Manage



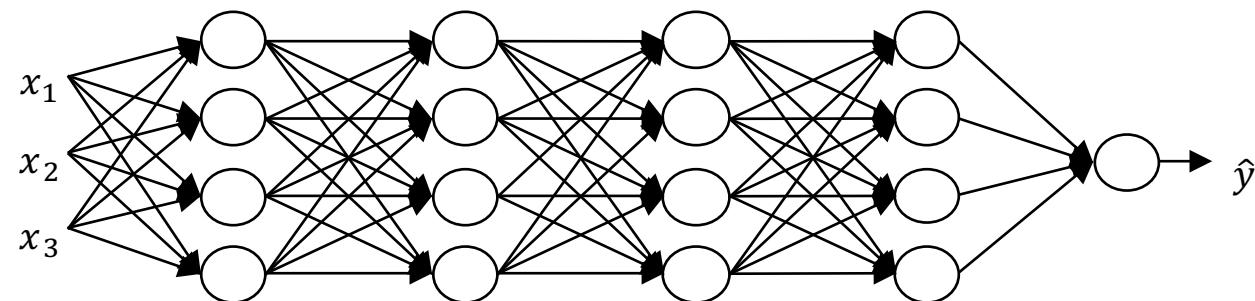
# Deep NN Training Focus

# Why Deep Learning?



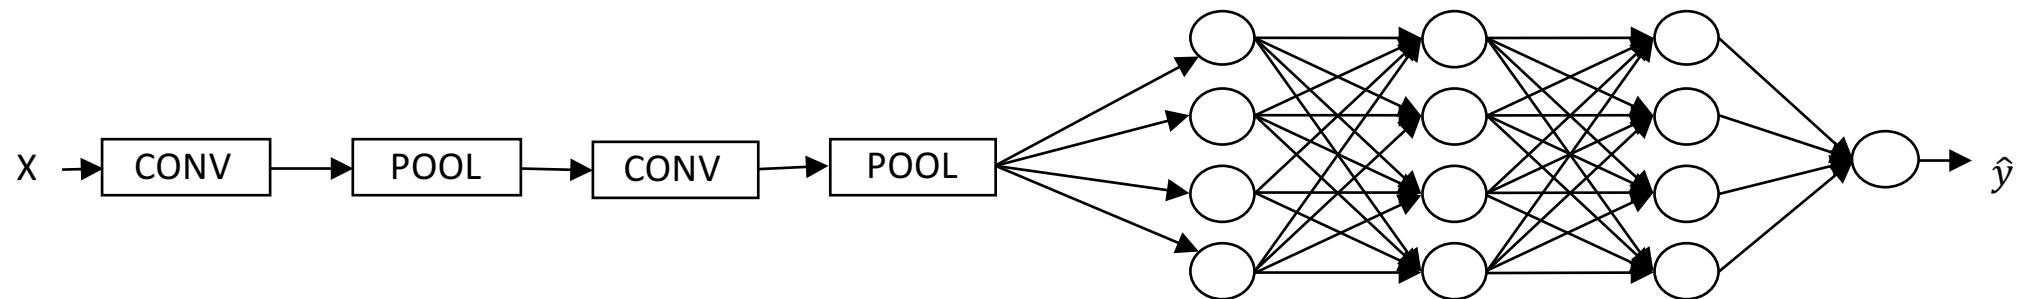
# Deep NN main architectures

- Feed Forward NN
  - One hidden layer
- Deep NN
  - Multiple hidden layers



# Deep NN main architectures

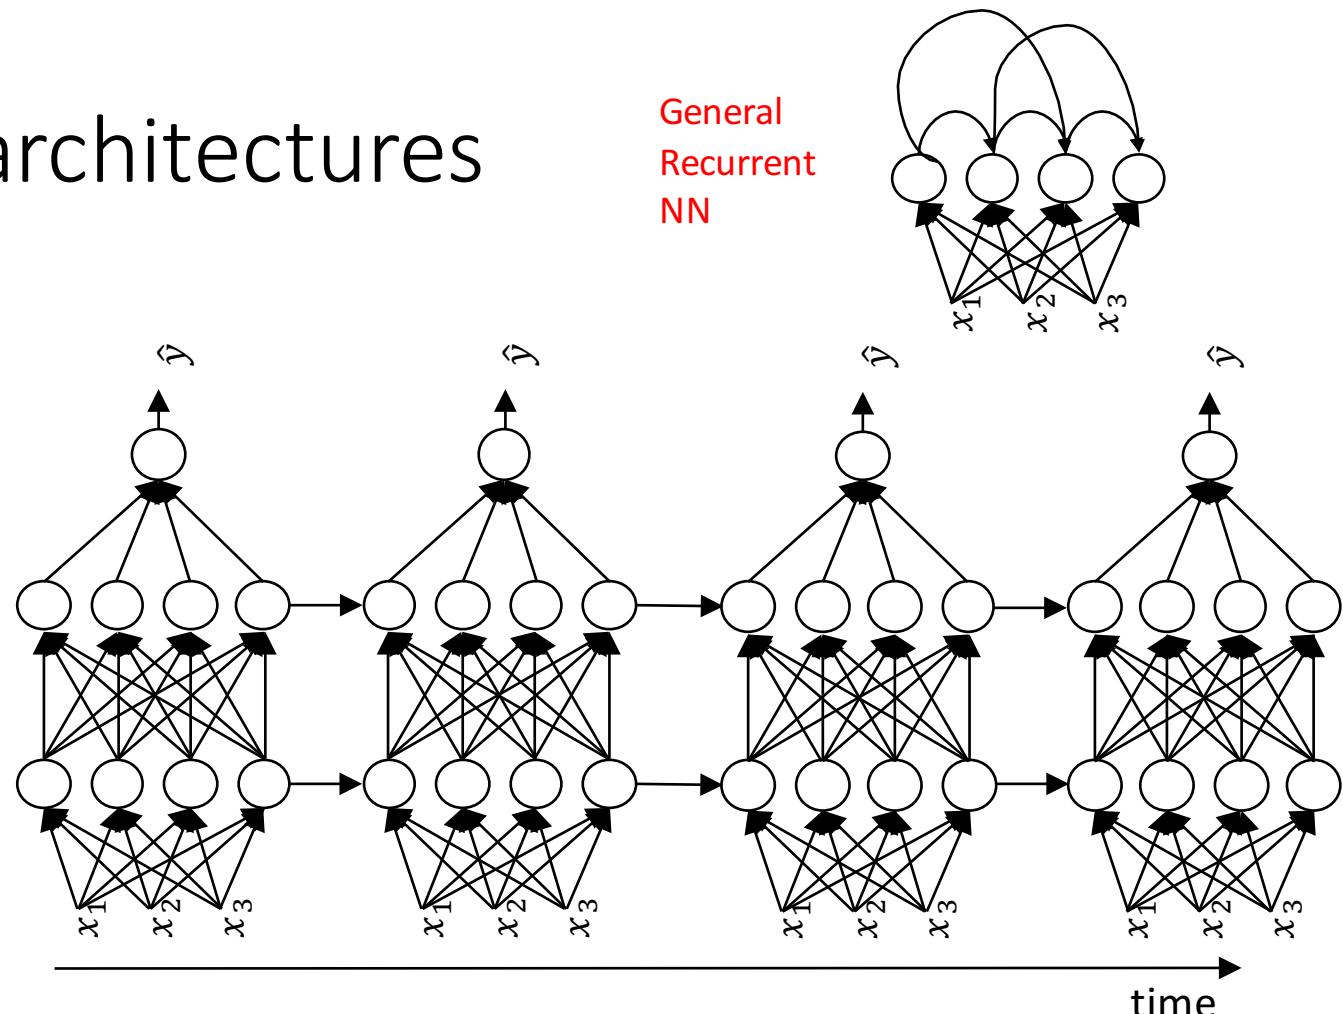
- **Convolutional NN**
  - Computer Vision
  - LeNet-5, AlexNet, VGG-16, ResNet, Inception



# Deep NN main architectures

- **Recurrent NN**

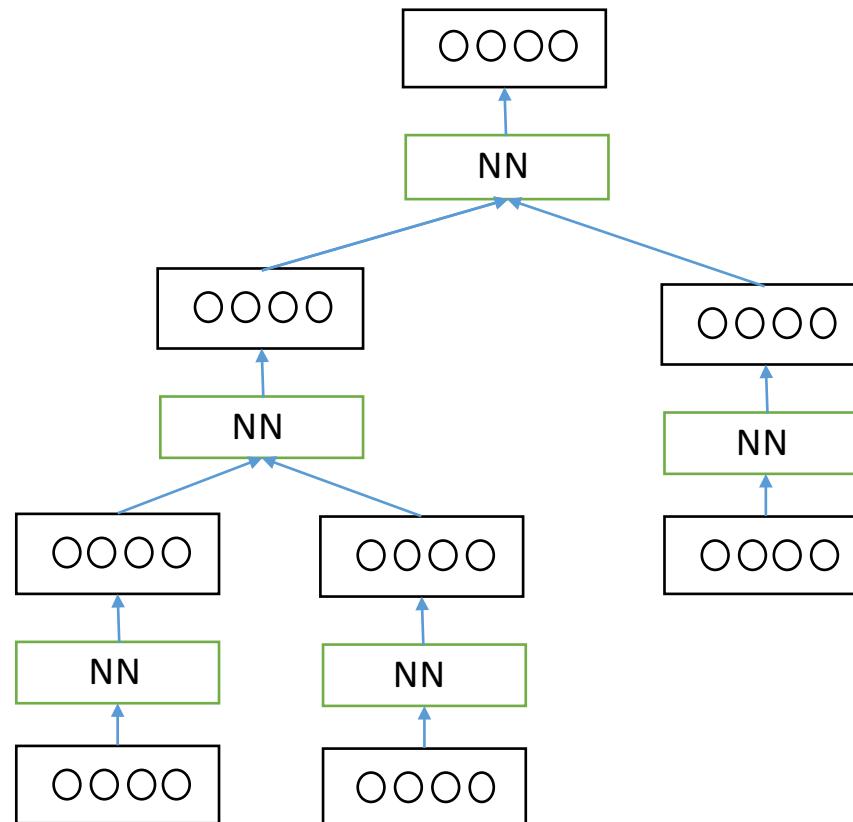
- Sequential models
  - NLP, Video, Signal
- Problem of vanishing gradient
- Gating Units: LSTM, GRU



# Deep NN main architectures

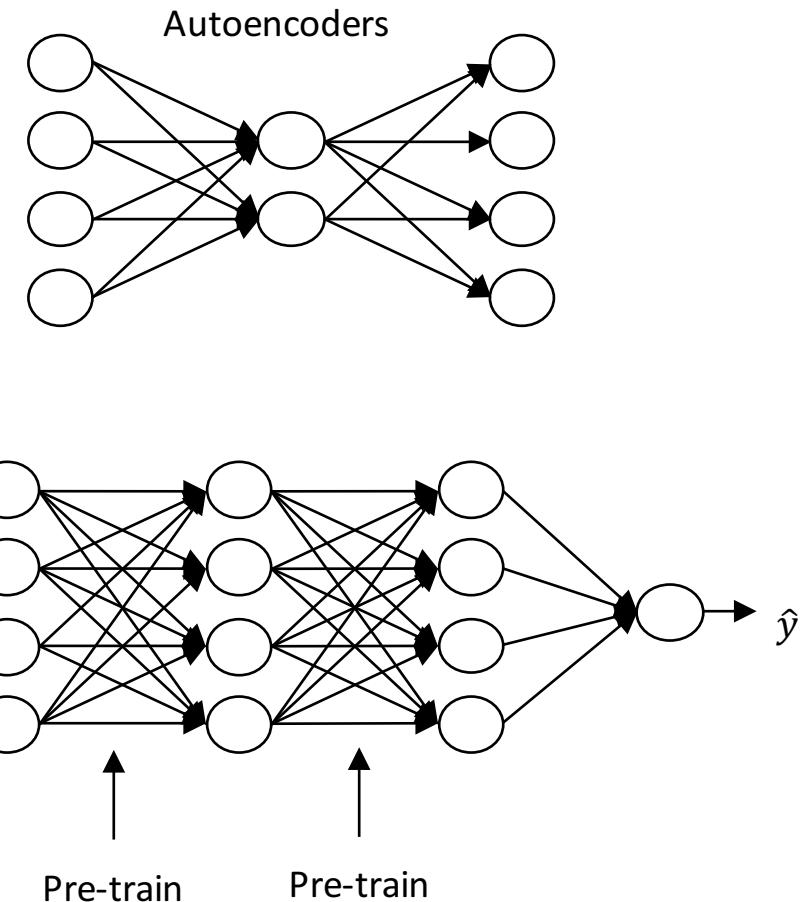
- Tree Recursive NN

- Understand compositionality
  - NLP parsing, Vision
  - Tree based
  - Recursive

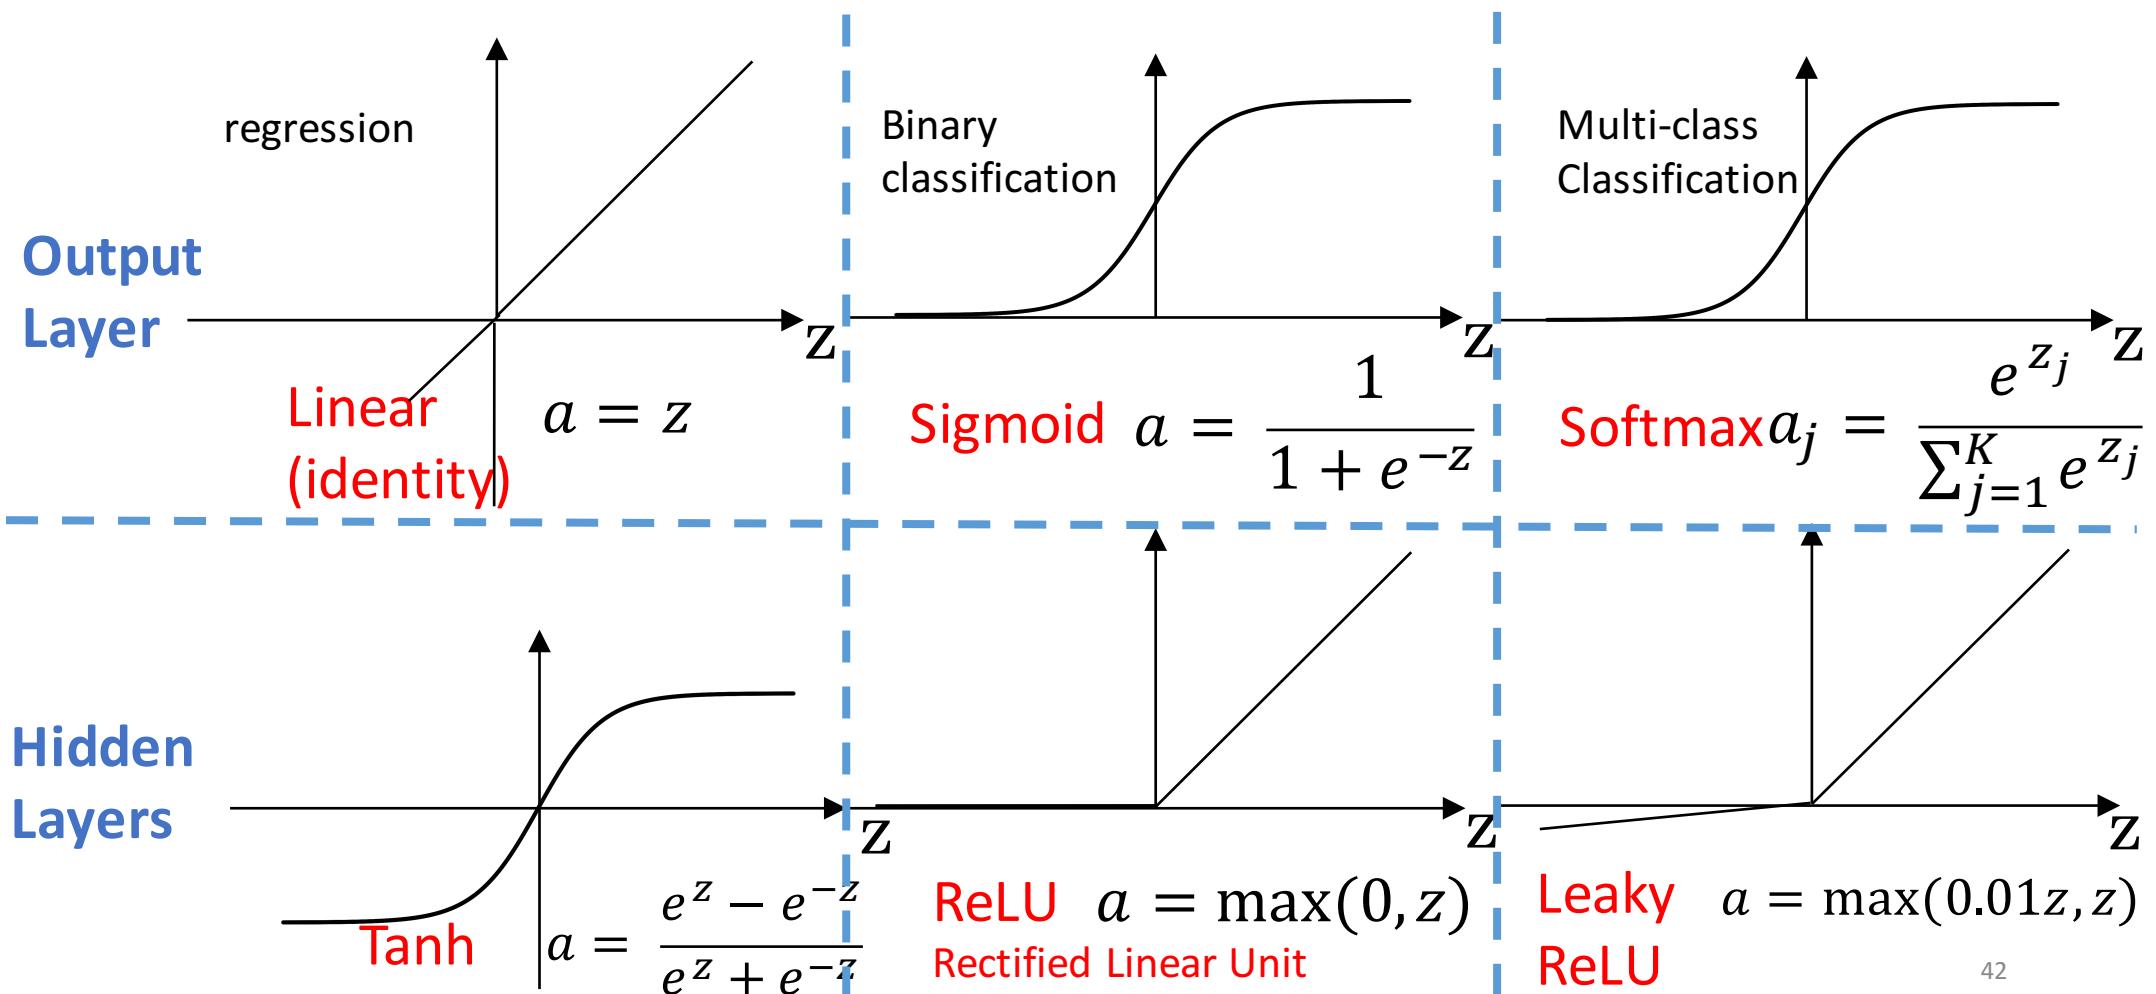


# Deep NN main architectures

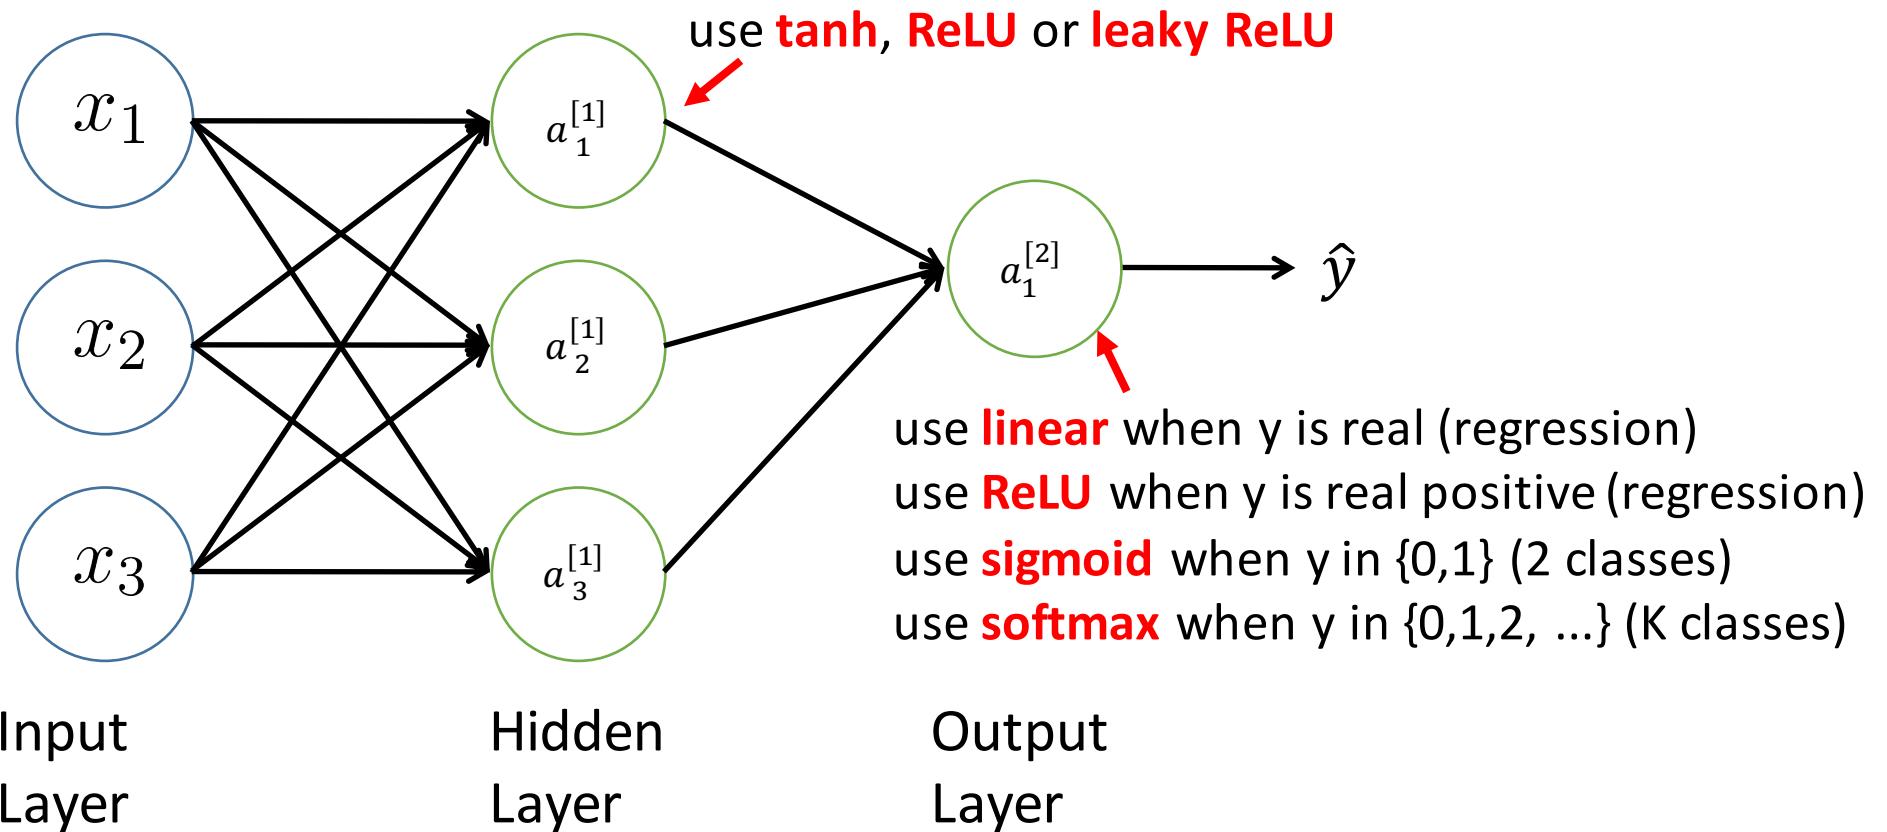
- **Unsupervised Pre-trained NN**
  - No Random Initialization
  - Parameters are initialized with unsupervised manner
    - Restricted Boltzman Machines
    - Autoencoders
  - Regularization effect
    - Solves Overfitting
  - Transfer Learning
    - Model trained on one task is re-purposed on a second related task



# Activation Functions

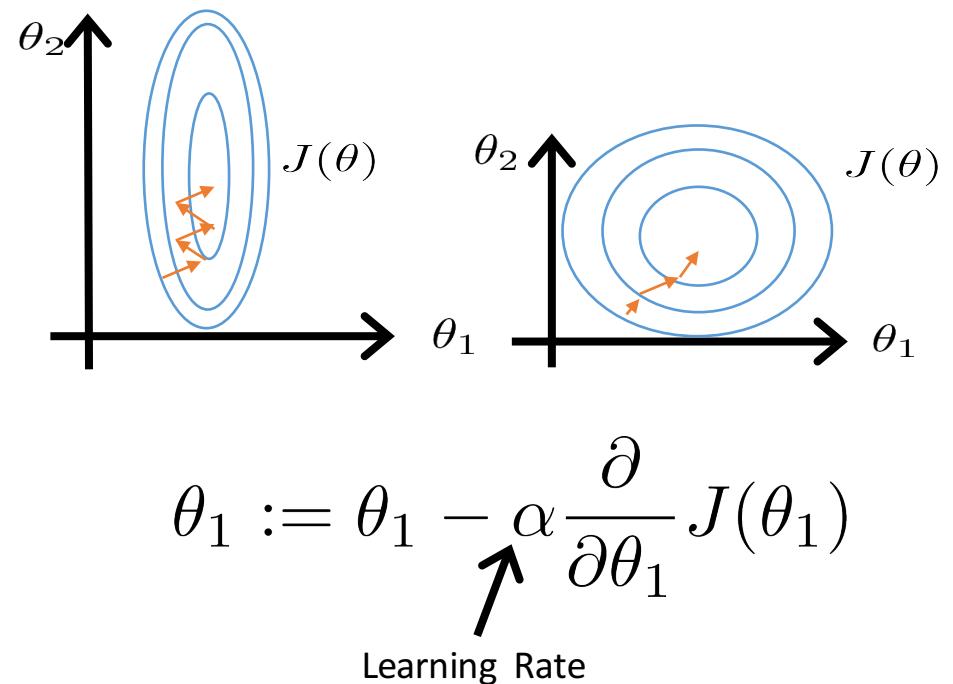


# Activation Functions



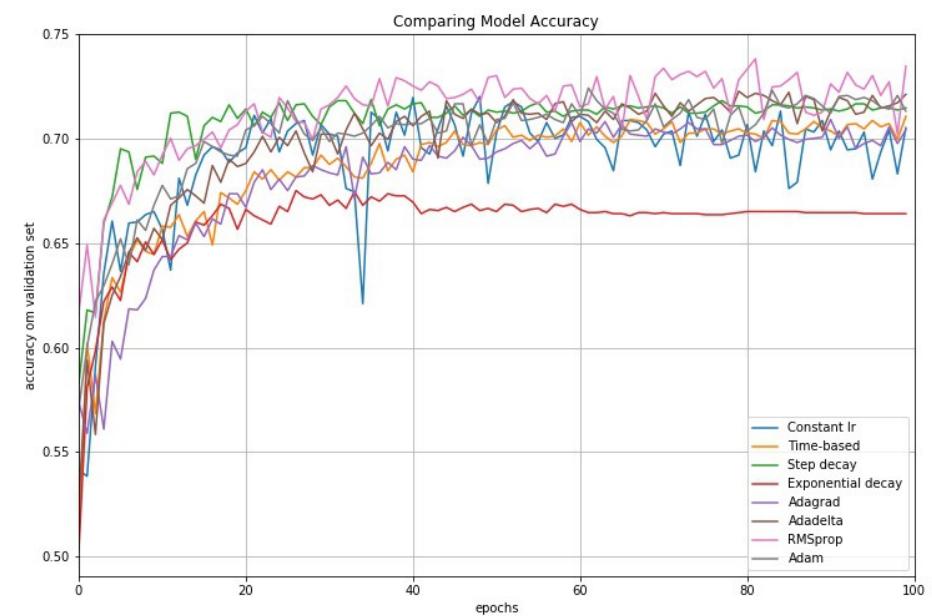
# Optimization (solve underfitting)

- **Back-Propagation**
- **Gradient Descent**
- **Normalizing Inputs** to speed up learning
- **Gradient checking** to prevent Vanishing/Exploding gradient
- **Momentum** to prevent oscillations



# Optimization (solve underfitting)

- **Gradient Descent** (Batch, Mini-batch, Stochastic)
  - Constant Learning Rate Decay
  - Time based decay:  $lr = lr_0 / (1 + kt)$
  - Step decay:  $lr = lr_0 * drop^{\lfloor epoch / epochs\_drop \rfloor}$
  - Exponential decay  $lr = lr_0 * \exp(-kt)$
- **Adaptive Learning Rate Decay**
  - Adagrad
  - Adadelta
  - RMSprop
  - Adam
- More details: [Yoshua Bengio Paper](#)



# Regularization (solve overfitting)

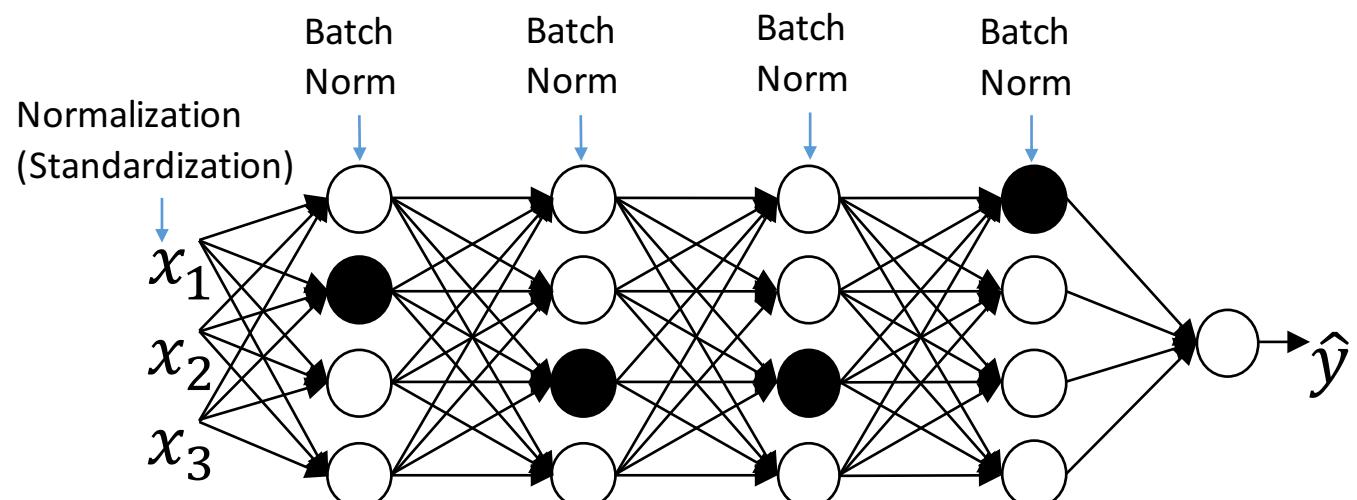
- Dropout
- Batch Normalization

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

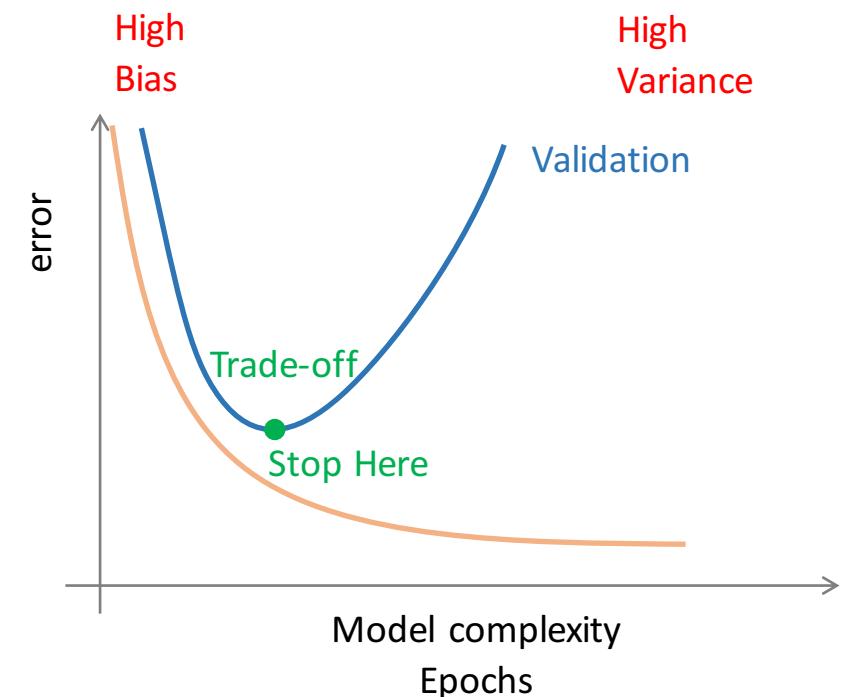
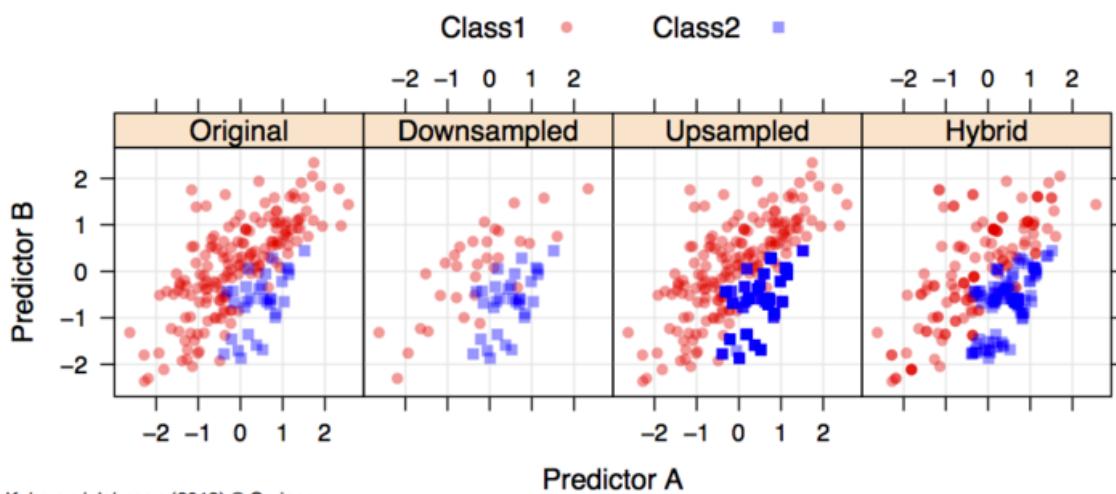
$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$$



# Regularization (solve overfitting)

- Data augmentation
- Early stopping



# How Can I Learn?

# How can I Learn?

- Math
  - Statistics, Probabilistic Graphical Models, Algebra, Optimization
- Programming Languages
  - Python, R, others
- Books
  - Gareth James, Daniela Witten, Trevor Hastie, Robert Tibshirani. “An introduction to statistical learning with applications in R”. 2013.
  - Tom M. Mitchell. “Machine Learning”. 1997
  - [Others](#)

# How can I Learn?

- MOOCs
  - Coursera: Machine learning, Deep learning specialization
  - Udacity, Udemy, Fun, etc.
- StackOverflow
- Research Papers
  - Read and rewrite algorithms from scratch
- Follow People
  - Andrew Ng, Yann LeCun, Jeff Hinton, Yoshua Bengio, Chris Manning, Sebastian Thrun, etc.

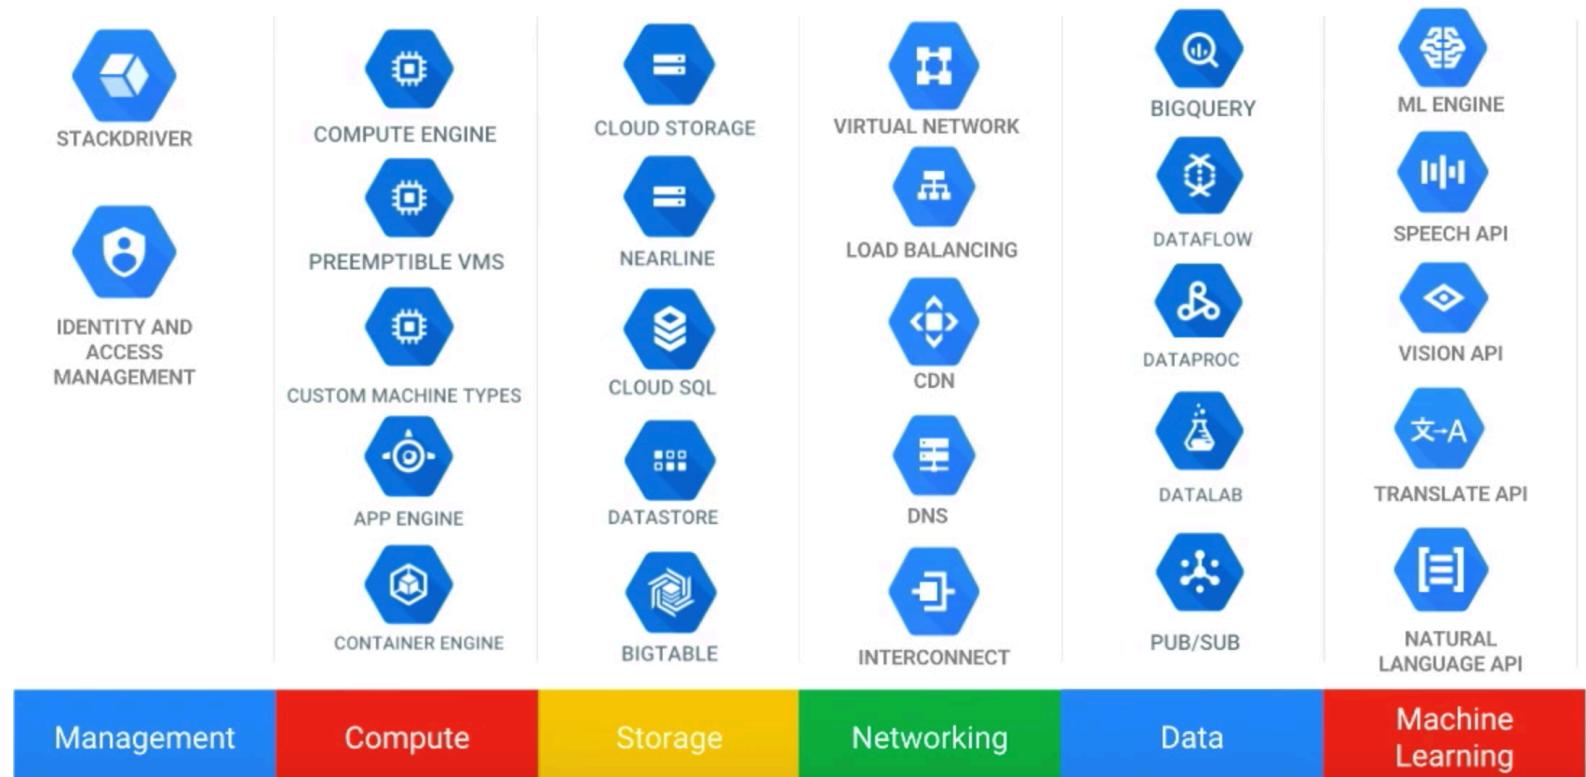
# How Can I Apply?

# How can I Apply?

- Start small projects and use Frameworks
  - [TensorFlow](#), [Keras](#), [Caffe](#), [Microsoft Cognitive Toolkit](#) (CNTK 2), [MXNet](#), [Scikit-learn](#), [Spark MLlib](#), etc.
- Kaggle, UCI Machine Learning Repo
  - Find data
  - Go for competitions
- Notebooks:
  - Jupyter notebook, google cloud Datalab
- Github
  - Find codes
  - Share your code
- Softwares
  - Knime: open source data analytics
  - IBM SPSS Modeler: data mining and text analytics software application from IBM

# How can I Apply?

## Google Cloud Services



# How can I Apply?

- Google ML APIs
  - Vision API
  - Video Intelligence API
  - Speech API
  - Natural Language API
  - Translation API



Cloud Translation API



Cloud Speech API



Cloud Vision API



Cloud Video Intelligence API



Cloud Natural Language API

# Demos

# My PhD students demos ([Youtube](#))

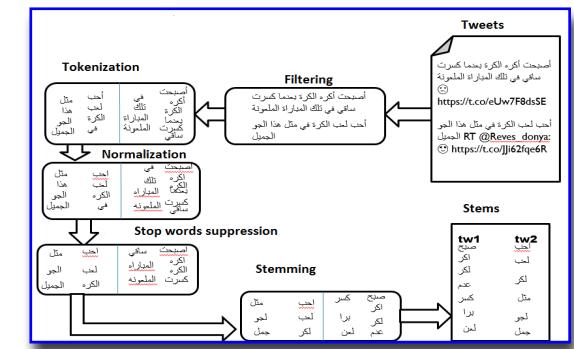
Younes Choubik



Marouane Hamda



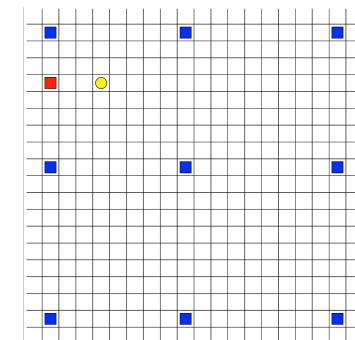
Hicham Boudkik



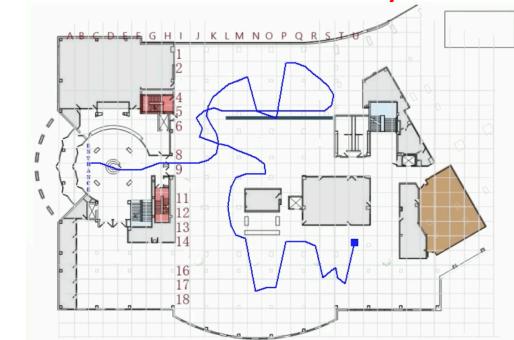
Marouane Benmoussa



Mohamed El Kaddoury



Mohamed Es-Salhy



THANKS!